Statistical theory of the continuous double auction

Eric Smith*, ¹ J. Doyne Farmer[†], ¹ László Gillemot, ¹ and Supriya Krishnamurthy ¹ Santa Fe Institute, 1399 Hyde Park Rd., Santa Fe NM 87501 (Dated: February 1, 2008)

Most modern financial markets use a continuous double auction mechanism to store and match orders and facilitate trading. In this paper we develop a microscopic dynamical statistical model for the continuous double auction under the assumption of IID random order flow, and analyze it using simulation, dimensional analysis, and theoretical tools based on mean field approximations. The model makes testable predictions for basic properties of markets, such as price volatility, the depth of stored supply and demand vs. price, the bid-ask spread, the price impact function, and the time and probability of filling orders. These predictions are based on properties of order flow and the limit order book, such as share volume of market and limit orders, cancellations, typical order size, and tick size. Because these quantities can all be measured directly there are no free parameters. We show that the order size, which can be cast as a nondimensional granularity parameter, is in most cases a more significant determinant of market behavior than tick size. We also provide an explanation for the observed highly concave nature of the price impact function. On a broader level, this work suggests how stochastic models based on zero-intelligence agents may be useful to probe the structure of market institutions. Like the model of perfect rationality, a stochastic-zero intelligence model can be used to make strong predictions based on a compact set of assumptions, even if these assumptions are not fully believable.

	Contents		3. Screening of the market-order rate	20
			4. Verifying the conservation laws	21
I.	Introduction	2	5. Self-consistent parametrization	21
	A. Motivation	2	6. Accounting for correlations	22
	B. Background: The continuous double auction	3	7. Generalizing the shift-induced source term	s 22
	C. The model	3	G. A mean-field theory of order separation	
	D. Summary of prior work	5	intervals: The Independent Interval	
тт		_	Approximation	24
11.	Overview of predictions of the model	5	1. Asymptotes and conservation rules	25
	A. Dimensional analysis	5	2. Direct simulation in interval coordinates	26
	B. Varying the granularity parameter ϵ	8		
	 Depth profile Liquidity for market orders: The price 	9	IV. Concluding remarks	28
	impact function	10	A. Ongoing work on empirical validation	28
	3. Spread	11	B. Future Enhancements	28
	4. Volatility and price diffusion	12	C. Comparison to standard models based on	20
	5. Liquidity for limit orders: Probability and		valuation and information arrival	29
	time to fill.	13		20
	C. Varying tick size dp/p_c	13	A. Relationship of Price impact to cumulative	Δ.
	2,2		depth	29
III.	Theoretical analysis	14	1. Moment expansion	30
	A. Summary of analytic methods	14	2. Quantiles	31
	B. Characterizing limit-order books: dual		2. Quantines	91
	coordinates	15	B. Supporting calculations in density	
	C. Frames and marginals	16	coordinates	32
	D. Factorization tests	17	1. Generating functional at general bin width	32
	E. Comments on renormalized diffusion	18	a. Recovering the continuum limit for prices	$\frac{32}{33}$
	F. Master equations and mean-field approximations	19	-	
	1. A number density master equation	19	2. Cataloging correlations	33
	2. Solution by generating functional	20	a. Getting the intercept right	34
	2. Solution by generating functional	20	b. Fokker-Planck expanding correlations	34
			Acknowledgments	35
	responding author: desmith@santafe.edu Kinsev Professor		References	35

I. INTRODUCTION

This section provides background and motivation, a description of the model, and some historical context for work in this area. Section II gives an overview of the phenomenology of the model, explaining how dimensional analysis applies in this context, and presenting a summary of numerical results. Section III develops an analytic treatment of model, explaining some of the numerical findings of Section II. We conclude in Section IV with a discussion of how the model may be enhanced to bring it closer to real-life markets, and some comments comparing the approach taken here to standard models based on information arrival and valuation.

A. Motivation

In this paper we analyze the continuous double auction trading mechanism under the assumption of random order flow, developing a model introduced in [1]. This analvsis produces quantitative predictions about the most basic properties of markets, such as volatility, depth of stored supply and demand, the bid-ask spread, the price impact, and probability and time to fill. These predictions are based on the rate at which orders flow into the market, and other parameters of the market, such as order size and tick size. The predictions are falsifiable with no free parameters. This extends the original random walk model of Bachelier [2] by providing a basis for the diffusion rate of prices. The model also provides a possible explanation for the highly concave nature of the price impact function. Even though some of the assumptions of the model are too simple to be literally true, the model provides a foundation onto which more realistic assumptions may easily be added.

The model demonstrates the importance of financial institutions in setting prices, and how solving a necessary economic function such as providing liquidity can have unanticipated side-effects. In a world of imperfect rationality and imperfect information, the task of demand storage necessarily causes persistence. Under perfect rationality all traders would instantly update their orders with the arrival of each piece of new information, but this is clearly not true for real markets. The limit order book, which is the queue used for storing unexecuted orders, has long memory when there are persistent orders. It can be regarded as a device for storing supply and demand, somewhat like a capacitor is a device for storing charge. We show that even under completely random IID order flow, the price process displays anomalous diffusion and interesting temporal structure. The converse is also interesting: For prices to be effectively random, incoming order flow must be non-random, in just the right way to compensate for the persistence. (See the remarks in Section IV C.)

This work is also of interest from a fundamental point of view because it suggests an alternative approach to doing economics. The assumption of perfect rationality has been popular in economics because it provides a parsimonious model that makes strong predictions. In the spirit of Gode and Sunder [3], we show that the opposite extreme of zero intelligence random behavior provides another reference model that also makes very strong predictions. Like perfect rationality, zero intelligence is an extreme simplification that is obviously not literally true. But as we show here, it provides a useful tool for probing the behavior of financial institutions. The resulting model may easily be extended by introducing simple boundedly rational behaviors. We also differ from standard treatments in that we do not attempt to understand the properties of prices from fundamental assumptions about utility. Rather, we split the problem in two. We attempt to understand how prices depend on order flow rates, leaving the problem of what determines these order flow rates for the future.

One of our main results concerns the average price impact function. The liquidity for executing a market order can be characterized by a price impact function $\Delta p = \phi(\omega, \tau, t)$. Δp is the shift in the logarithm of the price at time $t + \tau$ caused by a market order of size ω placed at time t. Understanding price impact is important for practical reasons such as minimizing transaction costs, and also because it is closely related to an excess demand function¹, providing a natural starting point for theories of statistical or dynamical properties of markets [4, 5]. A naive argument predicts that the price impact $\phi(\omega)$ should increase at least linearly. This argument goes as follows: Fractional price changes should not depend on the scale of price. Suppose buying a single share raises the price by a factor k > 1. If k is constant, buying ω shares in succession should raise it by k^{ω} . Thus, if buying ω shares all at once affects the price at least as much as buying them one at a time, the ratio of prices before and after impact should increase at least exponentially. Taking logarithms implies that the price impact as we have defined it above should increase at least linearly.²

In contrast, from empirical studies $\phi(\omega)$ for buy orders appears to be concave [6, 7, 8, 9, 10, 11]. Lillo *et al.* have shown for that for stocks in the NYSE the concave behavior of the price impact is quite consistent across different stocks [11]. Our model produces concave price impact functions that are in qualitative agreement with these results.

Our work also demonstrates the value of physics techniques for economic problems. Our analysis makes exten-

¹ In financial models it is common to define an excess demand function as demand minus supply; when the context is clear the modifier "excess" is dropped, so that demand refers to both supply and demand.

² This has practical implications. It is common practice to break up orders in order to reduce losses due to market impact. With a sufficiently concave market impact function, in contrast, it is cheaper to execute an order all at once.

sive use of dimensional analysis, the solution of a master equation through a generating functional, and a mean field approach that is commonly used to analyze non-equilibrium reaction-diffusion systems and evaporation-deposition problems.

B. Background: The continuous double auction

Most modern financial markets operate continuously. The mismatch between buyers and sellers that typically exists at any given instant is solved via an order-based market with two basic kinds of orders. Impatient traders submit market orders, which are requests to buy or sell a given number of shares immediately at the best available price. More patient traders submit *limit orders*, or quotes which also state a limit price, corresponding to the worst allowable price for the transaction. (Note that the word "quote" can be used either to refer to the limit price or to the limit order itself.) Limit orders often fail to result in an immediate transaction, and are stored in a queue called the *limit order book*. Buy limit orders are called bids, and sell limit orders are called offers or asks. We use the logarithmic price a(t) to denote the position of the best (lowest) offer and b(t) for the position the best (highest) bid. These are also called the *inside* quotes. There is typically a non-zero price gap between them, called the spread s(t) = a(t) - b(t). Prices are not continuous, but rather have discrete quanta called ticks. Throughout this paper, all prices will be expressed as logarithms, and to avoid endless repetition, the word price will mean the logarithm of the price. The minimum interval that prices change on is the tick size dp (also defined on a logarithmic scale; note this is not true for real markets). Note that dp is not necessarily infinitesimal.

As market orders arrive they are matched against limit orders of the opposite sign in order of first price and then arrival time, as shown in Fig. 1. Because orders are placed for varying numbers of shares, matching is not necessarily one-to-one. For example, suppose the best offer is for 200 shares at \$60 and the the next best is for 300 shares at \$60.25; a buy market order for 250 shares buys 200 shares at \$60 and 50 shares at \$60.25, moving the best offer a(t) from \$60 to \$60.25. A high density of limit orders per price results in high *liquidity* for market orders, i.e., it decreases the price movement when a market order is placed. Let n(p,t) be the stored density of limit order volume at price p, which we will call the depth profile of the limit order book at any given time t. The total stored limit order volume at price level pis n(p,t)dp. For unit order size the shift in the best ask a(t) produced by a buy market order is given by solving the equation

$$\omega = \sum_{p=a(t)}^{p'} n(p,t)dp \tag{1}$$

for p'. The shift in the best ask p' - a(t), where is the

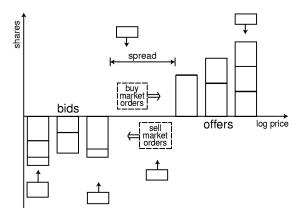


FIG. 1: A schematic illustration of the continuous double auction mechanism and our model of it. Limit orders are stored in the limit order book. We adopt the arbitrary convention that buy orders are negative and sell orders are positive. As a market order arrives, it has transactions with limit orders of the opposite sign, in order of price (first) and time of arrival (second). The best quotes at prices a(t) or b(t) move whenever an incoming market order has sufficient size to fully deplete the stored volume at a(t) or b(t). Our model assumes that market order arrival, limit order arrival, and limit order cancellation follow a Poisson process. New offers (sell limit orders) can be placed at any price greater than the best bid, and are shown here as "raining down" on the price axis. Similarly, new bids (buy limit orders) can be placed at any price less than the best offer. Bids and offers that fall inside the spread become the new best bids and offers. All prices in this model are logarithmic.

instantaneous price impact for buy market orders. A similar statement applies for sell market orders, where the price impact can be defined in terms of the shift in the best bid. (Alternatively, it is also possible to define the price impact in terms of the change in the midpoint price).

We will refer to a buy limit order whose limit price is greater than the best ask, or a sell limit order whose limit price is less than the best bid, as a *crossing limit order* or *marketable limit order*. Such limit orders result in immediate transactions, with at least part of the order immediately executed.

C. The model

This model introduced in reference [1], is designed to be as analytically tractable as possible while capturing key features of the continuous double auction. All the order flows are modeled as Poisson processes. We assume that market orders arrive in chunks of σ shares, at a rate of μ shares per unit time. The market order may be a 'buy' order or a 'sell' order with equal probability. (Thus the rate at which buy orders or sell orders arrive individually is $\mu/2$.) Limit orders arrive in chunks of σ shares as well, at a rate α shares per unit price and per

unit time for buy orders and also for sell orders. Offers are placed with uniform probability at integer multiples of a tick size dp in the range of price b(t) , andsimilarly for bids on $-\infty . When a market$ order arrives it causes a transaction; under the assumption of constant order size, a buy market order removes an offer at price a(t), and if it was the last offer at that price, moves the best ask up to the next occupied price tick. Similarly, a sell market order removes a bid at price b(t), and if it is the last bid at that price, moves the best bid down to the next occupied price tick. In addition, limit orders may also be removed spontaneously by being canceled or by expiring, even without a transaction having taken place. We model this by letting them be removed randomly with constant probability δ per unit time.

While the assumption of limit order placement over an infinite interval is clearly unrealistic, it provides a tractable boundary condition for modeling the behavior of the limit order book near the midpoint price m(t) = (a(t)+b(t))/2, which is the region of interest since it is where transactions occur. Limit orders far from the midpoint are usually canceled before they are executed (we demonstrate this later in Fig. 5), and so far from the midpoint, limit order arrival and cancellation have a steady state behavior characterized by a simple Poisson distribution. Although under the limit order placement process the total number of orders placed per unit time is infinite, the order placement per unit price interval is bounded and thus the assumption of an infinite interval creates no problems. Indeed, it guarantees that there are always an infinite number of limit orders of both signs stored in the book, so that the bid and ask are always well-defined and the book never empties. (Under other assumptions about limit order placement this is not necessarily true, as we later demonstrate in Fig. 30.) We are also considering versions of the model involving more realistic order placement functions; see the discussion in Section IVB.

In this model, to keep things simple, we are using the conceptual simplification of effective market orders and effective limit orders. When a crossing limit order is placed part of it may be executed immediately. The effect of this part on the price is indistinguishable from that of a market order of the same size. Similarly, given that this market order has been placed, the remaining part is equivalent to a non-crossing limit order of the same size. Thus a crossing limit order can be modeled as an effective market order followed by an effective (non-crossing) limit order.³ Working in terms of effective market and limit orders affects data analysis: The effective market order arrival rate μ combines both pure market orders

and the immediately executed components of crossing limit orders, and similarly the limit order arrival rate α corresponds only to the components of limit orders that are not executed immediately. This is consistent with the boundary conditions for the order placement process, since an offer with $p \leq b(t)$ or a bid with $p \geq a(t)$ would result in an immediate transaction, and thus would be effectively the same as a market order. Defining the order placement process with these boundary conditions realistically allows limit orders to be placed anywhere inside the spread.

Another simplification of this model is the use of logarithmic prices, both for the order placement process and for the tick size dp. This has the important advantage that it ensures that prices are always positive. In real markets price ticks are linear, and the use of logarithmic price ticks is an approximation that makes both the calculations and the simulation more convenient. We find that the limit $dp \to 0$, where tick size is irrelevant, is a good approximation for many purposes. We find that tick size is less important than other parameters of the problem, which provides some justification for the approximation of logarithmic price ticks.

Assuming a constant probability for cancellation is clearly $ad\ hoc$, but in simulations we find that other assumptions with well-defined timescales, such as constant duration time, give similar results. For our analytic model we use a constant order size σ . In simulations we also use variable order size, e.g. half-normal distributions with standard deviation $\sqrt{\pi/2\sigma}$, which ensures that the mean value remains σ . As long as these distributions have thin tails, the differences do not qualitatively affect most of the results reported here, except in a trivial way. As discussed in Section IVB, decay processes without well-defined characteristic times and size distributions with power law tails give qualitatively different results and will be treated elsewhere.

Even though this model is simply defined, the time evolution is not trivial. One can think of the dynamics as being composed of three parts: (1) the buy market order/sell limit order interaction, which determines the best ask: (2) the sell market order/buy limit order interaction, which determines the best bid; and (3) the random cancellation process. Processes (1) and (2) determine each others' boundary conditions. That is, process (1) determines the best ask, which sets the boundary condition for limit order placement in process (2), and process (2) determines the best bid, which determines the boundary conditions for limit order placement in process (1). Thus processes (1) and (2) are strongly coupled. It is this coupling that causes the bid and ask to remain close to each other, and guarantees that the spread s(t) = a(t) - b(t) is a stationary random variable, even though the bid and ask are not. It is the coupling of these processes through their boundary conditions that provides the nonlinear feedback that makes the price process complex.

³ In assigning independently random distributions for the two events, our model neglects the correlation between market and limit order arrival induced by crossing limit orders.

D. Summary of prior work

There are two independent lines of prior work, one in the financial economics literature, and the other in the physics literature. The models in the economics literature are directed toward empirical analysis, and treat the order process as static. In contrast, the models in the physics literature are conceptual toy models, but they allow the order process to react to changes in prices, and are thus fully dynamic. Our model bridges this gap. This is explained in more detail below.

The first model of this type that we are aware of was due to Mendelson [12], who modeled random order placement with periodic clearing. This was developed along different directions by Cohen et al. [13], who used techniques from queuing theory, but assumed only one price level and addressed the issue of time priority at that level (motivated by the existence of a specialist who effectively pinned prices to make them stationary). Domowitz and Wang [14] and Bollerslev et al. [15] further developed this to allow more general order placement processes that depend on prices, but without solving the full dynamical problem. This allows them to get a stationary solution for prices. In contrast, in our model the prices that emerge make a random walk, and so are much more realistic. In order to get a solution for the depth of the order book we have to go into price coordinates that comove with the random walk. Dealing with the feedback between order placement and prices makes the problem much more difficult, but it is key for getting reasonable results.

The models in the physics literature incorporate price dynamics, but have tended to be conceptual toy models designed to understand the anomalous diffusion properties of prices. This line of work begins with a paper by Bak et al. [16] which was developed by Eliezer and Kogan [17] and by Tang [18]. They assume that limit orders are placed at a fixed distance from the midpoint, and that the limit prices of these orders are then randomly shuffled until they result in transactions. It is the random shuffling that causes price diffusion. This assumption, which we feel is unrealistic, was made to take advantage of the analogy to a standard reaction-diffusion model in the physics literature. Maslov [19] introduced an alterative model that was solved analytically in the mean-field limit by Slanina [20]. Each order is randomly chosen to be either a buy or a sell, and either a limit order or a market order. If a limit order, it is randomly placed within a fixed distance of the current price. This again gives rise to anomalous price diffusion. A model allowing limit orders with Poisson order cancellation was proposed by Challet and Stinchcombe [21]. Iori and Chiarella [22] have numerically studied a model including fundamentalists and technical traders.

The model studied in this paper was introduced by Daniels *et al.* [1]. This adds to the literature by introducing a model that treats the feedback between order placement and price movement, while having enough re-

alism so that the parameters can be tested against real data. The prior models in the physics literature have tended to focus primarily on the anomalous diffusion of prices. While interesting and important for refining risk calculations, this is a second-order effect. In contrast, we focus on the first order effects of primary interest to market participants, such as the bid-ask spread, volatility, depth profile, price impact, and the probability and time to fill an order. We demonstrate how dimensional analysis becomes a useful tool in an economic setting, and develop mean field theories in a context that is more challenging than that of the toy models of previous work.

Subsequent to reference [1], Bouchaud et al. [23] demonstrated that, under the assumption that prices execute a random walk, by introducing an additional free parameter they can derive a simple equation for the depth profile. In this paper we show how to do this from first principles without introducing a free parameter.

II. OVERVIEW OF PREDICTIONS OF THE MODEL

In this section we give an overview of the phenomenology of the model. Because this model has five parameters, understanding all their effects would generally be a complicated problem in and of itself. This task is greatly simplified by the use of dimensional analysis, which reduces the number of independent parameters from five to two. Thus, before we can even review the results, we need to first explain how dimensional analysis applies in this setting. One of the surprising aspects of this model is that one can derive several powerful results using the simple technique of dimensional analysis alone.

Unless otherwise mentioned the results presented in this section are based on simulations. These results are compared to theoretical predictions in Section III.

A. Dimensional analysis

Because dimensional analysis is not commonly used in economics we first present a brief review. For more details see Bridgman [24].

Dimensional analysis is a technique that is commonly used in physics and engineering to reduce the number of independent degrees of freedom by taking advantage of the constraints imposed by dimensionality. For sufficiently constrained problems it can be used to guess the answer to a problem without doing a full analysis. The idea is to write down all the factors that a given phenomenon can depend on, and then find the combination that has the correct dimensions. For example, consider the problem of the period of a pendulum: The period T has dimensions of time. Obvious candidates that it might depend on are the mass of the bob m (which has units of mass), the length l (which has units of distance), and the acceleration of gravity q (which has units

Parameter	Description	Dimensions	
α	limit order rate	shares/(price	time)
μ	market order rate	shares/time	
δ	order cancellation rate	1/time	
dp	tick size	price	
σ	characteristic order size	shares	

TABLE I: The five parameters that characterize this model. α , μ , and δ are order flow rates, and dp and σ are discreteness parameters.

of $distance/time^2$). There is only one way to combine these to produce something with dimensions of time, i.e. $T \sim \sqrt{l/g}$. This determines the correct formula for the period of a pendulum up to a constant. Note that it makes it clear that the period does not depend on the mass, a result that is not obvious a priori. We were lucky in this problem because there were three parameters and three dimensions, with a unique combination of the parameters having the right dimensions; in general dimensional analysis can only be used to reduce the number of free parameters through the constraints imposed by their dimensions.

For this problem the three fundamental dimensions in the model are shares, price, and time. Note that by price, we mean the logarithm of price; as long as we are consistent, this does not create problems with the dimensional analysis. There are five parameters: three rate constants and two discreteness parameters. The order flow rates are μ , the market order arrival rate, with dimensions of shares per time; α , the limit order arrival rate per unit price, with dimensions of shares per price per time; and δ , the rate of limit order decays, with dimensions of 1/time. These play a role similar to rate constants in physical problems. The two discreteness parameters are the price tick size dp, with dimensions of price, and the order size σ , with dimensions of shares. This is summarized in table I.

Dimensional analysis can be used to reduce the number of relevant parameters. Because there are five parameters and three dimensions (price, shares, time), and because in this case the dimensionality of the parameters is sufficiently rich, the dimensional relationships reduce the degrees of freedom, so that all the properties of the limit-order book can be described by functions of two parameters. It is useful to construct these two parameters so that they are nondimensional.

We perform the dimensional reduction of the model by guessing that the effect of the order flow rates is primary to that of the discreteness parameters. This leads us to construct nondimensional units based on the order flow parameters alone, and take nondimensionalized versions of the discreteness parameters as the independent parameters whose effects remain to be understood. As we will see, this is justified by the fact that many of the properties of the model depend only weakly on the discreteness parameters. We can thus understand much of the richness of the phenomenology of the model through

Parameter	Description	Expression
N_c	characteristic number of shares	$\mu/2\delta$
p_c	characteristic price interval	$\mu/2\alpha$
t_c	characteristic time	$1/\delta$
dp/p_c	nondimensional tick size	$2\alpha dp/\mu$
ϵ	nondimensional order size	$2\delta\sigma/\mu$

TABLE II: Important characteristic scales and nondimensional quantities. We summarize the characteristic share size, price and times defined by the order flow rates, as well as the two nondimensional scale parameters dp/p_c and ϵ that characterize the effect of finite tick size and order size. Dimensional analysis makes it clear that all the properties of the limit order book can be characterized in terms of functions of these two parameters.

dimensional analysis alone.

There are three order flow rates and three fundamental dimensions. If we temporarily ignore the discreteness parameters, there are unique combinations of the order flow rates with units of shares, price, and time. These define a characteristic number of shares $N_c = \mu/2\delta$, a characteristic price interval $p_c = \mu/2\alpha$, and a characteristic timescale $t_c = 1/\delta$. This is summarized in table II. The factors of two occur because we have defined the market order rate for either a buy or a sell order to be $\mu/2$. We can thus express everything in the model in nondimensional terms by dividing by N_c , p_c , or t_c as appropriate, e.g. to measure shares in nondimensional units $\hat{N} = N/N_c$, or to measure price in nondimensional units $\hat{p} = p/p_c$.

The value of using nondimensional units is illustrated in Fig. 2. Fig. 2(a) shows the average depth profile for three different values of μ and δ with the other parameters held fixed. When we plot these results in dimensional units the results look quite different. However, when we plot them in terms of nondimensional units, as shown in Fig. 2(b), the results are indistinguishable. As explained below, because we have kept the nondimensional order size fixed, the collapse is perfect. Thus, the problem of understanding the behavior of this model is reduced to studying the effect of tick size and order size.

To understand the effect of tick size and order size it is useful to do so in nondimensional terms. The nondimensional scale parameter based on tick size is constructed by dividing by the characteristic price, i.e. $dp/p_c=2\alpha dp/\mu$. The theoretical analysis and the simulations show that there is a sensible continuum limit as the tick size $dp\to 0$, in the sense that there is non-zero price diffusion and a finite spread. Furthermore, the dependence on tick size is weak, and for many purposes the limit $dp\to 0$ approximates the case of finite tick size fairly well. As we will see, working in this limit is essential for getting tractable analytic results.

A nondimensional scale parameter based on order size is constructed by dividing the typical order size (which is measured in shares) by the characteristic number of shares N_c , i.e. $\epsilon \equiv \sigma/N_c = 2\delta\sigma/\mu$. ϵ characterizes the "chunkiness" of the orders stored in the limit order

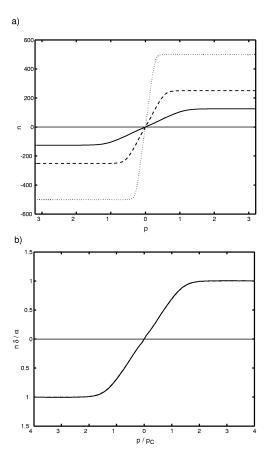


FIG. 2: The usefulness of nondimensional units. (a) We show the average depth profile for three different parameter sets. The parameters $\alpha=0.5,\ \sigma=1,\ \text{and}\ dp=0$ are held constant, while δ and μ are varied. The line types are: (dotted) $\delta=0.001,\ \mu=0.2;\ (\text{dashed})\ \delta=0.002,\ \mu=0.4$ and (solid) $\delta=0.004,\ \mu=0.8.$ (b) is the same, but plotted in nondimensional units. The horizontal axis has units of price, and so has nondimensional units $\hat{p}=p/pc=2\alpha p/\mu$. The vertical axis has units of $n\ shares/price$, and so has nondimensional units $\hat{n}=np_c/N_c=n\delta/\alpha$. Because we have chosen the parameters to keep the nondimensional order size ϵ constant, the collapse is perfect. Varying the tick size has little effect on the results other than making them discrete.

book. As we will see, ϵ is an important determinant of liquidity, and it is a particularly important determinant of volatility. In the continuum limit $\epsilon \to 0$ there is no price diffusion. This is because price diffusion can occur only if there is a finite probability for price levels outside the spread to be empty, thus allowing the best bid or ask to make a persistent shift. If we let $\epsilon \to 0$ while the average depth is held fixed the number of individual orders becomes infinite, and the probability that spontaneous decays or market orders can create gaps outside the spread becomes zero. This is verified in simulations. Thus the limit $\epsilon \to 0$ is always a poor approximation to a real market. ϵ is a more important parameter than the tick size dp/p_c . In the mean field analysis in Section III,

Quantity	Dimensions	Scaling relation
Asymptotic depth	shares/price	$d \sim \alpha/\delta$
Spread	price	$s \sim \mu/\alpha$
Slope of depth profile	$shares/price^2$	$\lambda \sim \alpha^2/\mu\delta = d/s$
Price diffusion rate	$price^2/time$	$D_0 \sim \mu^2 \delta / \alpha^2$

TABLE III: Estimates from dimensional analysis for the scaling of a few market properties based on order flow rates alone. α is the limit order density rate, μ is the market order rate, and δ is the spontaneous limit order removal rate. These estimates are constructed by taking the combinations of these three rates that have the proper units. They neglect the dependence on on the order granularity ϵ and the nondimensional tick size dp/p_c . More accurate relations from simulation and theory are given in table IV.

we let $dp/p_c \to 0$, reducing the number of independent parameters from two to one, and in many cases find that this is a good approximation.

The order size σ can be thought of as the order granularity. Just as the properties of a beach with fine sand are quite different from that of one populated by fist-sized boulders, a market with many small orders behaves quite differently from one with a few large orders. N_c provides the scale against which the order size is measured, and ϵ characterizes the granularity in relative terms. Alternatively, $1/\epsilon$ can be thought of as the annihilation rate from market orders expressed in units of the size of spontaneous decays. Note that in nondimensional units the number of shares can also be written $\hat{N} = N/N_c = N\epsilon/\sigma$.

The construction of the nondimensional granularity parameter illustrates the importance of including a spontaneous decay process in this model. If $\delta=0$ (which implies $\epsilon=0$) there is no spontaneous decay of orders, and depending on the relative values of μ and α , generically either the depth of orders will accumulate without bound or the spread will become infinite. As long as $\delta>0$, in contrast, this is not a problem.

For some purposes the effects of varying tick size and order size are fairly small, and we can derive approximate formulas using dimensional analysis based only on the order flow rates. For example, in table III we give dimensional scaling formulas for the average spread, the market order liquidity (as measured by the average slope of the depth profile near the midpoint), the volatility, and the asymptotic depth (defined below). Because these estimates neglect the effects of discreteness, they are only approximations of the true behavior of the model, which do a better job of explaining some properties than others. Our numerical and analytical results show that some quantities also depend on the granularity parameter ϵ and to a weaker extent on the tick size dp/p_c . Nonetheless, the dimensional estimates based on order flow alone provide a good starting point for understanding market behavior. A comparison to more precise formulas derived from theory and simulations is given in table IV.

An approximate formula for the mean spread can be derived by noting that it has dimensions of *price*, and the unique combination of order flow rates with these dimensions

Quantity	Scaling relation	Figure
Asymptotic depth	$d = \alpha/\delta$	3
Spread	$s = (\mu/\alpha)f(\epsilon, dp/p_c)$	10, 24
Slope of depth profile	$\lambda = (\alpha^2/\mu\delta)g(\epsilon, dp/p_c)$	3, 20 - 21
Price diffusion $(\tau \to 0)$	$D_0 = (\mu^2 \delta / \alpha^2) \epsilon^{-0.5}$	11, 14(c)
Price diffusion $(\tau \to \infty)$	$D_{\infty} = (\mu^2 \delta / \alpha^2) \epsilon^{0.5}$	11, 14(c)

TABLE IV: The dependence of market properties on model parameters based on simulation and theory, with the relevant figure numbers. These formulas include corrections for order granularity ϵ and finite tick size dp/p_c . The formula for asymptotic depth from dimensional analysis in table III is exact with zero tick size. The expression for the mean spread is modified by a function of ϵ and dp/p_c , though the dependence on them is fairly weak. For the liquidity λ , corresponding to the slope of the depth profile near the origin, the dimensional estimate must be modified because the depth profile is no longer linear (mainly depending on ϵ) and so the slope depends on price. The formulas for the volatility are empirical estimates from simulations. The dimensional estimate for the volatility from Table III is modified by a factor of $\epsilon^{-0.5}$ for the early time price diffusion rate and a factor of $\epsilon^{0.5}$ for the late time price diffusion rate.

sions is μ/α . While the dimensions indicate the scaling of the spread, they cannot determine multiplicative factors of order unity. A more intuitive argument can be made by noting that inside the spread removal due to cancellation is dominated by removal due to market orders. Thus the total limit order placement rate inside the spread, for either buy or sell limit orders αs , must equal the order removal rate $\mu/2$, which implies that spread is $s = \mu/2\alpha$. As we will see later, this argument can be generalized and made more precise within our mean-field analysis which then also predicts the observed dependence on the granularity parameter ϵ . However this dependence is rather weak and only causes a variation of roughly a factor of two for $\epsilon < 1$ (see Figs. 10 and 24), and the factor of 1/2derived above is a good first approximation. Note that this prediction of the mean spread is just the characteristic price p_c .

It is also easy to derive the mean asymptotic depth, which is the density of shares far away from the midpoint. The asymptotic depth is an artificial construct of our assumption of order placement over an infinite interval; it should be regarded as providing a simple boundary condition so that we can study the behavior near the midpoint price. The mean asymptotic depth has dimensions of shares/price, and is therefore given by α/δ . Furthermore, because removal by market orders is insignificant in this regime, it is determined by the balance between order placement and decay, and far from the midpoint the depth at any given price is Poisson distributed. This result is exact.

The average slope of the depth profile near the midpoint is an important determinant of liquidity, since it affects the expected price response when a market order arrives. The slope has dimensions of *shares/price*², which implies that in terms of the order flow rates it scales roughly as $\alpha^2/\mu\delta$. This is also the ratio of the asymptotic depth to the spread. As we will see later, this is a good approximation when $\epsilon \sim 0.01$, but for smaller values of ϵ the depth profile is not linear near the midpoint, and this approximation fails.

The last two entries in table IV are empirical estimates for the price diffusion rate D, which is proportional to the square of the volatility. That is, for normal diffusion, starting from a point at t=0, the variance v after time t is v = Dt. The volatility at any given timescale t is the square root of the variance at timescale t. The estimate for the diffusion rate based on dimensional analysis in terms of the order flow rates alone is $\mu^2 \delta/\alpha^2$. However, simulations show that short time diffusion is much faster than long time diffusion, due to negative autocorrelations in the price process, as shown in Fig. 11. The initial and the asymptotic diffusion rates appear to obey the scaling relationships given in table IV. Though our mean-field theory is not able to predict this functional form, the fact that early and late time diffusion rates are different can be understood within the framework of our analysis, as described in Sec. IIIE. Anomalous diffusion of this type implies negative autocorrelations in midpoint prices. Note that we use the term "anomalous diffusion" to imply that the diffusion rate is different on short and long timescales. We do not use this term in the sense that it is normally used in the physics literature, i.e. that the long-time diffusion is proportional to t^{γ} with $\gamma \neq 1$ (for long times $\gamma = 1$ in our case).

B. Varying the granularity parameter ϵ

We first investigate the effect of varying the order granularity ϵ in the limit $dp \to 0$. As we will see, the granularity has an important effect on most of the properties of the model, and particularly on depth, price impact, and price diffusion. The behavior can be divided into three regimes, roughly as follows:

- Large ϵ , i.e. $\epsilon \gtrsim 0.1$. This corresponds to a large accumulation of orders at the best bid and ask, nearly linear market impact, and roughly equal short and long time price diffusion rates. This is the regime where the mean-field approximation used in the theoretical analysis works best.
- Medium ϵ i.e. $\epsilon \sim 0.01$. In this range the accumulation of orders at the best bid and ask is small and near the midpoint price the depth profile increases nearly linearly with price. As a result, as a crude approximation the price impact increases as roughly the square root of order size.
- Small ϵ i.e. $\epsilon \lesssim 0.001$. The accumulation of orders at the best bid and ask is very small, and near the midpoint the depth profile is a convex function of price. The price impact is very concave. The short

time price diffusion rate is much greater than the long time price diffusion rate.

Since the results for bids are symmetric with those for offers about p=0, for convenience we only show the results for offers, i.e. buy market orders and sell limit orders. In this sub-section prices are measured relative to the midpoint, and simulations are in the continuum limit where the tick size $dp\to 0$. The results in this section are from numerical simulations. Also, bear in mind that far from the midpoint the predictions of this model are not valid due to the unrealistic assumption of an order placement process with an infinite domain. Thus the results are potentially relevant to real markets only when the price p is at most a few times as large as the characteristic price p_c .

1. Depth profile

The mean depth profile, i.e. the average number of shares per price interval, and the mean cumulative depth profile are shown in Fig. 3, and the standard deviation of the cumulative profile is shown in Fig. 4. Since the depth profile has units of shares/price, nondimensional units of depth profile are $\hat{n} = np_c/N_c = n\delta/\alpha$. The cumulative depth profile at any given time t is defined as

$$N(p,t) = \sum_{\tilde{p}=0}^{p} n(\tilde{p},t)dp. \tag{2}$$

This has units of shares and so in nondimensional terms is $\hat{N}(p) = N(p)/N_c = 2\delta N(p)/\mu = N(p)\epsilon/\sigma$.

In the high ϵ regime the annihilation rate due to market orders is low (relative to $\delta\sigma$), and there is a significant accumulation of orders at the best ask, so that the average depth is much greater than zero at the midpoint. The mean depth profile is a concave function of price. In the medium ϵ regime the market order removal rate increases, depleting the average depth near the best ask, and the profile is nearly linear over the range $p/p_c \leq 1$. In the small ϵ regime the market order removal rate increases even further, making the average depth near the ask very close to zero, and the profile is a convex function over the range $p/p_c \leq 1$.

The standard deviation of the depth profile is shown in Fig. 4. We see that the standard deviation of the cumulative depth is comparable to the mean depth, and that as ϵ increases, near the midpoint there is a similar transition from convex to concave behavior.

The uniform order placement process seems at first glance one of the most unrealistic assumptions of our model, leading to depth profiles with a finite asymptotic depth (which also implies that there are an infinite number of orders in the book). However, orders far away from the spread in the asymptotic region almost never get executed and thus do not affect the market dynamics. To demonstrate this in Fig. 5 we show the comparison between the limit-order depth profile and the depth

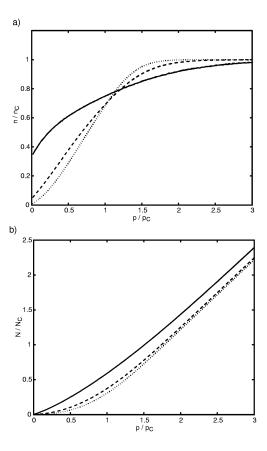


FIG. 3: The mean depth profile and cumulative depth versus $\hat{p}=p/p_c=2\alpha p/\mu$. The origin $p/p_c=0$ corresponds to the midpoint. (a) is the average depth profile n in nondimensional coordinates $\hat{n}=np_c/N_c=n\delta/\alpha$. (b) is nondimensional cumulative depth $N(p)/N_c$. We show three different values of the nondimensional granularity parameter: $\epsilon=0.2$ (solid), $\epsilon=0.02$ (dash), $\epsilon=0.02$ (dot), all with tick size dp=0.

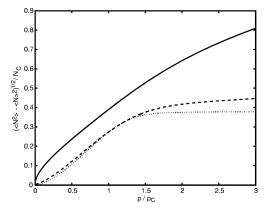


FIG. 4: Standard deviation of the nondimensionalized cumulative depth versus nondimensional price, corresponding to Fig. (3).

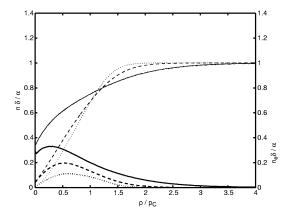


FIG. 5: A comparison between the depth profiles and the effective depth profiles as defined in the text, for different values of ϵ . Heavy lines refer to the effective depth profiles n_e and the light lines correspond to the depth profiles.

 n_e of only those orders which eventually get executed. The density n_e of executed orders decreases rapidly as a function of the distance from the mid-price. Therefore we expect that near the midpoint our results should be similar to alternative order placement processes, as long as they also lead to an exponentially decaying profile of executed orders (which is what we observe above). However, to understand the behavior further away from the midpoint we are also working on enhancements that include more realistic order placement processes grounded on empirical measurements of market data, as summarized in section IV B.

2. Liquidity for market orders: The price impact function

In this sub-section we study the instantaneous price impact function $\phi(t,\omega,\tau\to 0)$. This is defined as the (logarithm of the) midpoint price shift immediately after the arrival of a market order in the absence of any other events. This should be distinguished from the asymptotic price impact $\phi(t,\omega,\tau\to\infty)$, which describes the permanent price shift. While the permanent price shift is clearly very important, we do not study it here. The reader should bear in mind that all prices p, a(t), etc. are logarithmic.

The price impact function provides a measure of the liquidity for executing market orders. (The liquidity for limit orders, in contrast, is given by the probability of execution, studied in section IIB5). At any given time t, the instantaneous ($\tau = 0$) price impact function is the

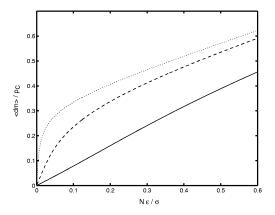


FIG. 6: The average price impact corresponding to the results in Fig. (3). The average instantaneous movement of the nondimensional mid-price, $\langle dm \rangle/p_c$ caused by an order of size $N/N_c = N\epsilon/\sigma$. $\epsilon = 0.2$ (solid), $\epsilon = 0.02$ (dash), $\epsilon = 0.002$ (dot).

inverse of the cumulative depth profile. This follows immediately from equations (1) and (2), which in the limit $dp \to 0$ can be replaced by the continuum transaction equation:

$$\omega = N(p,t) = \int_0^p n(\tilde{p},t)d\tilde{p} \tag{3}$$

This equation makes it clear that at any fixed t the price impact can be regarded as the inverse of the cumulative depth profile N(p,t). When the fluctuations are sufficiently small we can replace n(p,t) by its mean value $n(p) = \langle n(p,t) \rangle$. In general, however, the fluctuations can be large, and the average of the inverse is not equal to the inverse of the average. There are corrections based on higher order moments of the depth profile, as given in the moment expansion derived in Appendix A1. Nonetheless, the inverse of the mean cumulative depth provides a qualitative approximation that gives insight into the behavior of the price impact function. (Note that everything becomes much simpler using medians, since the median of the cumulative price impact function is exactly the inverse of the median price impact, as derived in Appendix A 1).

Mean price impact functions are shown in Fig. 6 and the standard deviation of the price impact is shown in Fig. 7. The price impact exhibits very large fluctuations for all values of ϵ : The standard deviation has the same order of magnitude as the mean or even greater for small $N\epsilon/\sigma$ values. Note that these are actually virtual price impact functions. That is, to explore the behavior of the instantaneous price impact for a wide range of order sizes, we periodically compute the price impact that an order of a given size would have caused at that instant, if it had been submitted. We have checked that real price impact curves are the same, but they require a much longer time to accumulate reasonable statistics.

⁴ Note that the ratio n_e/n is not the same as the probability of filling orders (Fig. 12) because in that case the price p/p_c refers to the distance of the order from the midpoint at the time when it was placed.

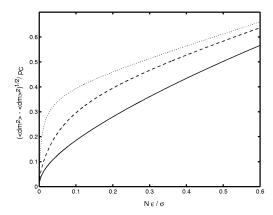


FIG. 7: The standard deviation of the instantaneous price impact dm/p_c corresponding to the means in Fig. 6, as a function of normalized order size $\epsilon N/\sigma$. $\epsilon = 0.2$ (solid), $\epsilon = 0.02$ (dash), $\epsilon = 0.002$ (dot).

One of the interesting results in Fig. 6 is the scale of the price impact. The price impact is measured relative to the characteristic price scale p_c , which as we have mentioned earlier is roughly equal to the mean spread. As we will argue in relation to Fig. 8, the range of nondimensional shares shown on the horizontal axis spans the range of reasonable order sizes. This figure demonstrates that throughout this range the price is the order of magnitude (and typically less than) the mean spread size.

Due to the accumulation of orders at the ask in the large ϵ regime, for small p the mean price impact is roughly linear. This follows from equation (3) under the assumption that n(p) is constant. In the medium ϵ regime, under the assumption that the variance in depth can be neglected, the mean price impact should increase as roughly $\omega^{1/2}$. This follows from equation (3) under the assumption that n(p) is linearly increasing and $n(0) \approx 0$. (Note that we see this as a crude approximation, but there can be substantial corrections caused by the variance of the depth profile). Finally, in the small ϵ regime the price impact is highly concave, increasing much slower than $\omega^{1/2}$. This follows because $n(0) \approx 0$ and the depth profile n(p) is convex.

To get a better feel for the functional form of the price impact function, in Fig. 8 we numerically differentiate it versus log order size, and plot the result as a function of the appropriately scaled order size. (Note that because our prices are logarithmic, the vertical axis already incorporates the logarithm). If we were to fit a local power law approximation to the function at each price, this corresponds to the exponent of that power law near that price. Notice that the exponent is almost always less than one, so that the price impact is almost always concave. Making the assumption that the effect of the variance of the depth is not too large, so that equation (3) is a good assumption, the behavior of this figure can be understood as follows: For $N/N_c \approx 0$ the price impact is dominated

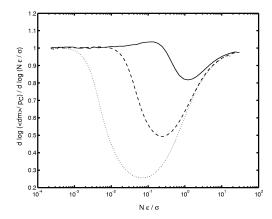


FIG. 8: Derivative of the nondimensional mean mid-price movement, with respect to logarithm of the nondimensional order size $N/N_c = N\epsilon/\sigma$, obtained from the price impact curves in Fig. 6.

by n(0) (the constant term in the average depth profile) and so the logarithmic slope of the price impact is always near to one. As N/N_c increases, the logarithmic slope is driven by the shape of the average depth profile, which is linear or convex for smaller ϵ , resulting in concave price impact. For large values of N/N_c , we reach the asymptotic region where the depth profile is flat (and where our model is invalid by design). Of course, there can be deviations to this behavior caused by the fact that the mean of the inverse depth profile is not in general the inverse of the mean, i.e. $\langle N^{-1}(p) \rangle \neq \langle N(p) \rangle^{-1}$ (see App. A 1).

To compare to real data, note that $N/N_c = N\epsilon/\sigma$. N/σ is just the order size in shares in relation to the average order size, so by definition it has a typical value of one. For the London Stock Exchange, we have found that typical values of ϵ are in the range 0.001-0.1. For a typical range of order sizes from 100-100,000 shares, with an average size of 10,000 shares, the meaningful range for N/N_c is therefore roughly 10^{-5} to 1. In this range, for small values of ϵ the exponent can reach values as low as 0.2. This offers a possible explanation for the previously mysterious concave nature of the price impact function, and contradicts the linear increase in price impact based on the naive argument presented in the introduction.

3. Spread

The probability density of the spread is shown in Fig. 9. This shows that the probability density is substantial at $s/p_c = 0$. (Remember that this is in the limit $dp \to 0$). The probability density reaches a maximum at a value of the spread approximately $0.2p_c$, and then decays. It might seem surprising at first that it decays more slowly for large ϵ , where there is a large accumulation of orders at the ask. However, it should be borne in mind

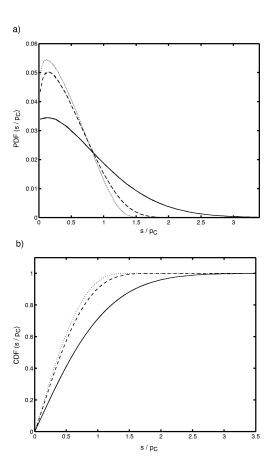


FIG. 9: The probability density function (a), and cumulative distribution function (b) of the nondimensionalized bid-ask spread s/p_c , corresponding to the results in Fig. (3). $\epsilon = 0.2$ (solid), $\epsilon = 0.02$ (dash), $\epsilon = 0.002$ (dot).

that the characteristic price $p_c = \mu/\alpha$ depends on ϵ . Since $\epsilon = 2\delta\sigma/\mu$, by eliminating μ this can be written $p_c = 2\sigma\delta/(\alpha\epsilon)$. Thus, holding the other parameters fixed, large ϵ corresponds to small p_c , and vice versa. So in fact, the spread is very small for large ϵ , and large for small ϵ , as expected. The figure just shows the small corrections to the large effects predicted by the dimensional scaling relations.

For large ϵ the probability density of the spread decays roughly exponentially moving away from the midpoint. This is because for large ϵ the fluctuations around the mean depth are roughly independent. Thus the probability for a market order to penetrate to a given price level is roughly the probability that all the ticks smaller than this price level contain no orders, which gives rise to an exponential decay. This is no longer true for small ϵ . Note that for small ϵ the probability distribution of the spread becomes insensitive to ϵ , i.e. the nondimensionalized distribution for $\epsilon=0.02$ is nearly the same as that for $\epsilon=0.002$.

It is apparent from Fig. 9 that in nondimensional units the mean spread increases with ϵ . This is confirmed in Fig. 10, which displays the mean value of the spread as a

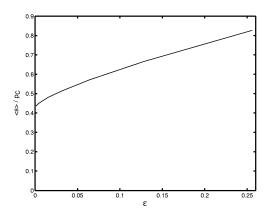


FIG. 10: The mean value of the spread in nondimensional units $\hat{s} = s/p_c$ as a function of ϵ . This demonstrates that the spread only depends weakly on ϵ , indicating that the prediction from dimensional analysis given in table (III) is a reasonable approximation.

function of ϵ . The mean spread increases monotonically with ϵ . It depends on ϵ as roughly a constant (equal to approximately 0.45 in nondimensional coordinates) plus a linear term whose slope is rather small. We believe that for most financial instruments $\epsilon < 0.3$. Thus the variation in the spread caused by varying ϵ in the range $0 < \epsilon < 0.3$ is not large, and the dimensional analysis based only on rate parameters given in table IV is a good approximation. We get an accurate prediction of the ϵ dependence across the full range of ϵ from the Independent Interval Approximation technique derived in section III G, as shown in Fig. 24.

4. Volatility and price diffusion

The price diffusion rate, which is proportional to the square of the volatility, is important for determining risk and is a property of central interest. From dimensional analysis in terms of the order flow rates the price diffusion rate has units of price²/time, and so must scale as $\mu^2 \delta / \alpha^2$. We can also make a crude argument for this as follows: The dimensional estimate of the spread (see Table IV) is $\mu/2\alpha$. Let this be the characteristic step size of a random walk, and let the step frequency be the characteristic time $1/\delta$ (which is the average lifetime for a share to be canceled). This argument also gives the above estimate for the diffusion rate. However, this is not correct in the presence of negative autocorrelations in the step sizes. The numerical results make it clear that there are important ϵ -dependent corrections to this result, as demonstrated below.

In Fig. 11 we plot simulation results for the variance of the change in the midpoint price at timescale τ , $\operatorname{Var}(m(t+\tau)-m(t))$. The slope is the diffusion rate, which at any fixed timescale is proportional to the square

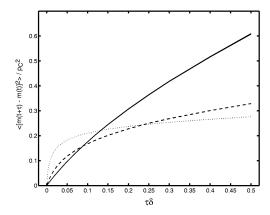


FIG. 11: The variance of the change in the nondimensionalized midpoint price versus the nondimensional time delay interval $\tau\delta$. For a pure random walk this would be a straight line whose slope is the diffusion rate, which is proportional to the square of the volatility. The fact that the slope is steeper for short times comes from the nontrivial temporal persistence of the order book. The three cases correspond to Fig. 3: $\epsilon = 0.2$ (solid), $\epsilon = 0.02$ (dash), $\epsilon = 0.002$ (dot).

of the volatility. It appears that there are at least two timescales involved, with a faster diffusion rate for short timescales and a slower diffusion rate for long timescales. Such anomalous diffusion is not predicted by mean-field analysis. Simulation results show that the diffusion rate is correctly described by the product of the estimate from dimensional analysis based on order flow parameters alone, $\mu^2 \delta/\alpha^2$, and a τ -dependent power of the nondimensional granularity parameter $\epsilon = 2\delta\sigma/\mu$, as summarized in table IV. We cannot currently explain why this power is -1/2 for short term diffusion and 1/2 for long-term diffusion. However, a qualitative understanding can be gained based on the conservation law we derive in Section III C. A discussion of how this relates to price diffusion is given in Section III E.

Note that the temporal structure in the diffusion process also implies non-zero autocorrelations of the midpoint price m(t). This corresponds to weak negative autocorrelations in price differences m(t)-m(t-1) that persist for timescales until the variance vs. τ becomes a straight line. The timescale depends on parameters, but is typically the order of 50 market order arrival times. This temporal structure implies that there exists an arbitrage opportunity which, when exploited, would make prices more random and the structure of the order flow non-random.

5. Liquidity for limit orders: Probability and time to fill.

The liquidity for limit orders depends on the probability that they will be filled, and the time to be filled. This obviously depends on price: Limit orders close to

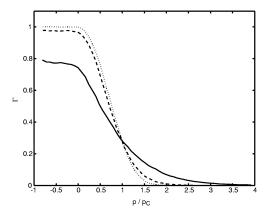


FIG. 12: The probability Γ for filling a limit order placed at a price p/p_c where p is calculated from the instantaneous midprice at the time of placement. The three cases correspond to Fig. 3: $\epsilon = 0.2$ (solid), $\epsilon = 0.02$ (dash), $\epsilon = 0.002$ (dot).

the current transaction prices are more likely to be filled quickly, while those far away have a lower likelihood to be filled. Fig. 12 plots the probability Γ of a limit order being filled versus the nondimensionalized price at which it was placed (as with all the figures in this section, this is shown in the midpoint-price centered frame). Fig. 12 shows that in nondimensional coordinates the probability of filling close to the bid for sell limit orders (or the ask for buy limit orders) decreases as ϵ increases. For large ϵ , this is less than 1 even for negative prices. This says that even for sell orders that are placed close to the best bid there is a significant chance that the offer is deleted before being executed. This is not true for smaller values of ϵ , where $\Gamma(0) \approx 1$. Far away from the spread the fill probabilities as a function of ϵ are reversed, i.e. the probability for filling limit orders increases as ϵ increases. The crossover point where the fill probabilities are roughly the same occurs at $p \approx p_c$. This is consistent with the depth profile in Fig. 3 which also shows that depth profiles for different values of ϵ cross at about $p \sim p_c$.

Similarly Fig 13 shows the average time τ taken to fill an order placed at a distance p from the instantaneous mid-price. Again we see that though the average time is larger at larger values of ϵ for small p/p_c , this behaviour reverses at $p \sim p_c$.

C. Varying tick size dp/p_c

The dependence on discrete tick size dp/p_c , of the cumulative distribution function for the spread, instantaneous price impact, and mid-price diffusion, are shown in Fig. 14. We chose an unrealistically large value of the tick size, with $dp/p_c=1$, to show that, even with very coarse ticks, the qualitative changes in behavior are typically relatively minor.

Fig. 14(a) shows the cumulative density function of

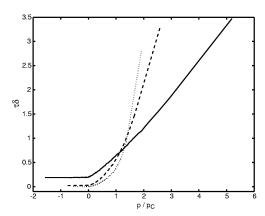


FIG. 13: The average time τ nondimensionalized by the rate δ , to fill a limit order placed at a distance p/p_c from the instantaneous mid-price.

the spread, comparing $dp/p_c=0$ and $dp/p_c=1$. It is apparent from this figure that the spread distribution for coarse ticks "effectively integrates" the distribution in the limit $dp \to 0$. That is, at integer tick values the mean cumulative depth profiles roughly match, and in between integer tick values, for coarse ticks the probability is smaller. This happens for the obvious reason that coarse ticks quantize the possible values of the spread, and place a lower limit of one tick on the value the spread can take. The shift in the mean spread from this effect is not shown, but it is consistent with this result; there is a constant offset of roughly 1/2 tick.

The alteration in the price impact is shown in Fig. 14(b). Unlike the spread distribution, the average price impact varies continuously. Even though the tick size is quantized, we are averaging over many events and the probability of a price impact of each tick size is a continuous function of the order size. Large tick size consistently lowers the price impact. The price impact rises more slowly for small p, but is then similar except for a downward translation.

The effect of coarse ticks is less trivial for mid-price diffusion, as shown in Fig. 14(c). At $\epsilon=0.002$, coarse ticks remove most of the rapid short-term volatility of the midpoint, which in the continuous-price case arises from price fluctuations smaller than $dp/p_c=1$. This lessens the negative autocorrelation of midpoint price returns, and reduces the anomalous diffusion. At $\epsilon=0.2$, where both early volatility and late negative autocorrelation are smaller, coarse ticks have less effect. The net result is that the mid-price diffusion becomes less sensitive to the value of ϵ as tick size increases, and there is less anomalous price diffusion.

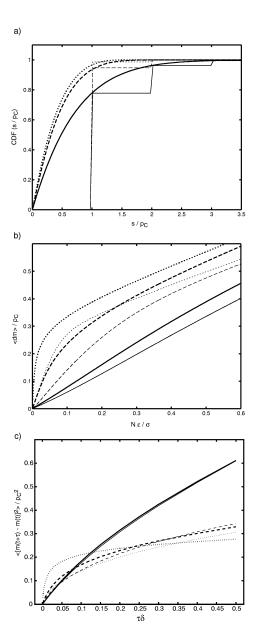


FIG. 14: Dependence of market properties on tick size. Heavy lines are $dp/p_c \rightarrow 0$; light lines are $dp/p_c = 1$. Cases correspond to Fig. 3, with $\epsilon = 0.2$ (solid), $\epsilon = 0.02$ (dash), $\epsilon = 0.002$ (dot). (a) is the cumulative distribution function for the nondimensionalized spread. (b) is instantaneous nondimensionalized price impact, (c) is diffusion of the nondimensionalized midpoint shift, corresponding to Fig. 11.

III. THEORETICAL ANALYSIS

A. Summary of analytic methods

We have investigated this model analytically using two approaches. The first one is based on a master equation, given in Section III F. This approach works best in the midpoint centered frame. Here we attempt to solve directly for the average number of shares at each price tick

as a function of price. The midpoint price makes a random walk with a nonstationary distribution. Thus the key to finding a stationary analytic solution for the average depth is to use comoving price coordinates, which are centered on a reference point near the center of the book, such as the midpoint or the best bid. In the first approximation, fluctuations about the mean depth at adjacent prices are treated as independent. This allows us to replace the distribution over depth profiles with a simpler probability density over occupation numbers n at each pand t. We can take a continuum limit by letting the tick size dp become infinitesimal. With finite order flow rates, this gives vanishing probability for the existence of more than one order at any tick as $dp \to 0$. This is described in detail in section IIIF 3. With this approach we are able to test the relevance of correlations as a function of the parameter ϵ as well as predict the functional dependence of the cumulative distribution of the spread on the depth profile. It is seen that correlations are negligible for large values of $\epsilon(\epsilon \sim 0.2)$ while they are very important for small values ($\epsilon \sim 0.002$).

Our second analytic approach which we term the Independent Interval Approximation (IIA) is most easily carried out in the bid-centered frame and is described in section III G. This approach uses a different representation, in which the solution is expressed in terms of the empty intervals between non-empty price ticks. The system is characterized at any instant of time by a set of intervals $\{...x_{-1}, x_0, x_1, x_2...\}$ where for example x_0 is the distance between the bid and the ask (the spread), x_{-1} is the distance between the second buy limit order and the bid and so on (see Fig. 15). Equations are written for how a given interval varies in time. Changes to adjacent intervals are related, giving us an infinite set of coupled non-linear equations. However using a meanfield approximation we are able to solve the equations. albeit only numerically. Besides predicting how the various intervals (for example the spread) vary with the parameters, this approach also predicts the depth profiles as a function of the parameters. The predictions from the IIA are compared to data from numerical simulations, in Section III G 2. They match very well for large ϵ and less well for smaller values of ϵ . The IIA can also be modified to incorporate various extensions to the model, as mentioned in Section III G 2.

In both approaches, we use a mean field approximation to get a solution. The approximation basically lies in assuming that fluctuations in adjacent intervals (which might be adjacent price ranges in the master equation approach or adjacent empty intervals in the IIA) are independent. Also, both approaches are most easily tractable only in the continuum limit $dp \to 0$, when every tick has at most only one order. They may however be extended to general tick size as well. This is explained in the appendix for the Master Equation approach.

Because correlations are important for small ϵ , both methods work well mostly in the large ϵ limit, though qualitative aspects of small ϵ behavior may also be

gleaned from them. Unfortunately, at least based on our preliminary investigation of London Stock Exchange data, it seems that it is this small ϵ limit that real markets may tend more towards. So our approximate solutions may not be as useful as we would like. Nonetheless, they do provide some conceptual insights into what determines depth and price impact.

In particular, we find that the shape of the mean depth profile depends on a single parameter ϵ , and that the relative sizes of its first few derivatives account for both the order size-dependence of the market impact, and the renormalization of the midpoint diffusivity. A higher relative rate of market versus limit orders depletes the center of the book, though less than the classical estimate predicts. This leads to more concave impact (explaining Fig. 8) and faster short-term diffusivity. However, the orders pile up more quickly (versus classically nondimensionalized price) with distance from the midpoint, causing the rapid early diffusion to suffer larger mean reversion. These are the effects shown in Fig. 11. We will elaborate on the above remarks in the following sections, however, the qualitative relation of impact to midpoint autocorrelation supplies a potential interpretation of data, which may be more robust than details of the model assumptions or its quantitative results.

Both of the treatments described above are approximations. We can derive an exact global conservation law of order placement and removal whose consequences we elaborate in section III C. This conservation law must be respected in any sensible analysis of the model, giving us a check on the approximations. It also provides some insight into the anomalous diffusion properties of this model.

B. Characterizing limit-order books: dual coordinates

We begin with the assumption of a *price space*. Price is a dimensional quantity, and the space is divided into bins of length dp representing the ticks, which may be finite or infinitesimal. Prices are then discrete or continuous-valued, respectively.

Statistical properties of interest are computed from temporal sequences or ensembles of limit-order book configurations. If n is the variable used to denote the number of shares from limit orders in some bin (p, p + dp) at the beginning t of an elementary time interval, a configuration is specified by a function n(p,t). It is convenient to take n positive for sell limit orders, and negative for buy limit orders. Because the model dynamics precludes crossing limit orders, there is in general a highest instantaneous buy limit-order price, called the bid b(t), and a lowest sell limit-order price, the ask a(t), with b(t) < a(t) always. The midpoint price, defined as $m(t) \equiv [a(t) + b(t)]/2$, may or may not be the price of any actual bin, if prices are discrete (m(t)) may be a half-integer multiple of dp). These quantities are diagrammed

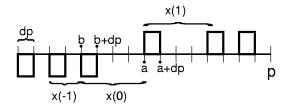


FIG. 15: The price space and order profile. n(p,t) has been chosen to be 0 or ± 1 , a restriction that will be convenient later. Price bins are labeled by their lower boundary price, and intervals x(N) will be defined below.

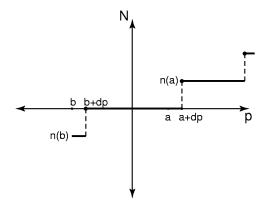


FIG. 16: The accumulated order number $N\left(p,t\right)$. $N\left(a,t\right)\equiv0$, because contributions from all bins cancel in the two sums. N remains zero down to $b\left(t\right)+dp$, because there are no uncanceled, nonzero terms. $N\left(b,t\right)$ becomes negative, because the second sum in Eq. (4) now contains $n\left(b,t\right)$, not canceled by the first.

in Fig. 15.

An equivalent specification of a limit-order book configuration is given by the cumulative order count

$$N\left(p,t\right) \equiv \sum_{-\infty}^{p-dp} \left| n\left(p,t\right) \right| - \sum_{-\infty}^{a-dp} \left| n\left(p,t\right) \right|, \tag{4}$$

where $-\infty$ denotes the lower boundary of the price space, whose exact value must not affect the results. (Because by definition there are no orders between the bid and ask, the bid could equivalently have been used as the origin of summation. Because price bins will be indexed here by their lower boundaries, though, it is convenient here to use the ask.) The absolute values have been placed so that N, like n, is negative in the range of buy orders and positive in the range of sells. The construction of N(p,t) is diagrammed in Fig. 16.

In many cases of either sparse orders or infinitesimal dp, with fixed order size (which we may as well define to be one share) there will be either zero or one share in any single bin, and Eq. (4) will be invertible to an equivalent specification of the limit-order book configuration

$$p(N,t) \equiv \max \{ p \mid N(p,t) = N \}, \tag{5}$$

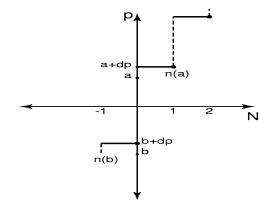


FIG. 17: The inverse function $p\left(N,t\right)$. The function is in general defined only on discrete values of N, so this domain is only invariant when order size is fixed, a convenience that will be assumed below. Between the discrete domain, and the definition of p as a maximum, the inverse function effectively interpolates between vertices of the reflected image of $N\left(p,t\right)$, as shown by the dotted line.

shown in Fig. 17. (Strictly, the inversion may be performed for any distribution of order sizes, but the resulting function is intrinsically discrete, so its domain is only invariant when order size is fixed. To give $p\left(N,t\right)$ the convenient properties of a well-defined function on an invariant domain, this will be assumed below.)

With definition (5), $p(0,t) \equiv a(t)$, $p(-1,t) \equiv b(t)$, and one can define the intervals between orders as

$$x(N,t) \equiv p(N,t) - p(N-1,t). \tag{6}$$

Thus x(0,t)=a(t)-b(t), the instantaneous bidask spread. The lowest values of x(N,t) bracketing the spread are shown in Fig. 15. For symmetric order-placement rules, probability distributions over configurations will be symmetric under either $n(p,t) \rightarrow -n(-p,t)$, or $x(N,t) \rightarrow x(-N,t)$. Coordinates N and p furnish a dual description of configurations, and n and x are their associated differences. The Master Equation approach of section III F assumes independent fluctuation in n while the Independent Interval Approximation of Sec. III G assumes independent fluctuation in x (In this section, it will be convenient to abbreviate $x(N,t) \equiv x_N(t)$).

C. Frames and marginals

The x(N,t) specification of limit-order book configurations has the property that its distribution is stationary under the dynamics considered here. The same is not true for p(N,t) or n(p,t) directly, because bid, midpoint, and ask prices undergo a random walk, with a renormalized diffusion coefficient. Stationary distributions for n-variables can be obtained in co-moving frames, of which there are several natural choices.

The bid-centered configuration is defined as

$$n_b(p,t) \equiv n(p-b(t),t). \tag{7}$$

If an appropriate rounding convention is adopted in the case of discrete prices, a *midpoint-centered configuration* can also be defined, as

$$n_{m}(p,t) \equiv n(p - m(t), t). \tag{8}$$

The midpoint-centered configuration has qualitative differences from the bid-centered configuration, which will be explored below. Both give useful insights to the order distribution and diffusion processes. The ask-centered configuration, $n_a(p,t)$, need not be considered if order placement and removal are symmetric, because it is a mirror image of $n_b(p,t)$.

The *spread* is defined as the difference $s(t) \equiv a(t) - b(t)$, and is the value of the ask in bid-centered coordinates. In midpoint-centered coordinates, the ask appears at s(t)/2.

The configurations n_b and n_m are dynamically correlated over short time intervals, but evolve ergodically in periods longer than finite characteristic correlation times. Marginal probability distributions for these can therefore be computed as time averages, either as functions on the whole price space, or at discrete sets of prices. Their marginal mean values at a single price p will be denoted $\langle n_b(p) \rangle$, $\langle n_m(p) \rangle$, respectively.

These means are subject to global balance constraints, between total order placement and removal in the price space. Because all limit orders are placed above the bid, the bid-centered configuration obeys a simple balance relation:

$$\frac{\mu}{2} = \sum_{p=b+dp}^{\infty} \left(\alpha - \delta \left\langle n_b \left(p \right) \right\rangle \right). \tag{9}$$

Eq. (9) says that buy market orders must account, on average, for the difference between all limit orders placed, and all decays. After passing to nondimensional coordinates below, this will imply an inverse relation between corrections to the classical estimate for diffusivity at early and late times, discussed in Sec. III E. In addition, this conservation law plays an important role in the analysis and determination of the x(N,t)'s, as we will see later in the text.

The midpoint-centered averages satisfy a different constraint:

$$\frac{\mu}{2} = \alpha \frac{\langle s \rangle}{2} + \sum_{p=b+dp}^{\infty} (\alpha - \delta \langle n_m(p) \rangle).$$
 (10)

Market orders in Eq. (10) account not only for the excess of limit order placement over evaporation at prices above the midpoint, but also the "excess" orders placed between b(t) and m(t). Since these always lead to midpoint shifts, they ultimately appear at positive comoving coordinates, altering the shape of $\langle n_m(p) \rangle$ relative to $\langle n_b(p) \rangle$. Their rate of arrival is $\alpha \langle m-b \rangle = \alpha \langle s \rangle /2$. These results are also confirmed in simulations.

D. Factorization tests

Whether in the bid-centered frame or the midpoint centered frame, the probability distribution function for the entire configuration $n\left(p\right)$ is too difficult a problem to solve in its entirety. However, an approximate master equation can be formed for n independently at each p if all joint probabilities factor into independent marginals, as

$$\Pr\left(\left\{n\left(p_{i}\right)\right\}_{i}\right) = \prod_{i} \Pr\left(n\left(p_{i}\right)\right), \tag{11}$$

where Pr denotes, for instance, a probability density for n orders in some interval around p.

Whenever orders are sufficiently sparse that the expected number in any price bin is simply the probability that the bin is occupied (up to a constant of proportionality), the independence assumption implies a relation between the cumulative distribution for the spread of the ask and the mean density profile. In units where the order size is one, the relation is

$$\Pr(s/2 < p) = 1 - \exp\left(-\sum_{p'=b+dp}^{p-dp} \langle n_m(p')\rangle\right). \quad (12)$$

This relation is tested against simulation results in Fig. 18. One can observe that there are three regimes.

A high- ϵ regime is defined when the mean density profile at the midpoint $\langle n_m(0)\rangle\lesssim 1$, and strongly concave downward. In this regime, the approximation of independent fluctuations is excellent, and a master equation treatment is expected to be useful. Intermediate- ϵ is defined by $\langle n_m(0)\rangle\ll 1$ and nearly linear, and the approximation of independence is marginal. Large- ϵ is defined by $\langle n_m(0)\rangle\ll 1$ and concave upward, and the approximation of independent fluctuations is completely invalid. These regimes of validity correspond also to the qualitative ranges noted already in Sec. II B.

In the bid centered frame however, Eq. 12 never seems to be valid for any range of parameters. We will discuss later why this might be so. For the present therefore, the master equation approach is carried out in the midpoint-centered frame. Alternatively, the mean field theory of the separations is most convenient in the bid-centered frame, so that frame will be studied in the dual basis. The relation of results in the two frames, and via the two methods of treatment, will provide a good qualitative, and for some properties quantitative, understanding of the depth profile and its effect on impacts.

It is possible in a modified treatment, to match certain features of simulations at any ϵ , by limited incorporation of correlated fluctuations. However, the general master equation will be developed independent of these, and tested against simulation results at large ϵ , where its defining assumptions are well met.

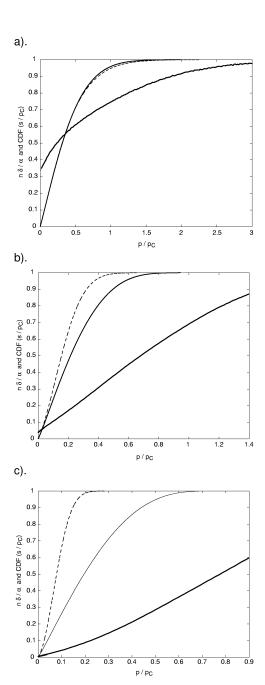


FIG. 18: CDFs $\Pr(s/2 < p)$ from simulations (thin solid), mean density profile $\langle n_m(p) \rangle$ from simulations (thick solid), and computed CDF of spread (thin dashed) from $\langle n_m(p) \rangle$, under the assumption of uncorrelated fluctuations, at three values of ϵ . (a): $\epsilon = 0.2$ (low market order rate); approximation is very good. (b): $\epsilon = 0.02$ (intermediate market order rate); approximation is marginal. (c): $\epsilon = 0.002$ (high market order rate); approximation is very poor.

E. Comments on renormalized diffusion

A qualitative understanding of why the diffusivity is different over short and long times scales, as well as why it may depend on ϵ , may be gleaned from the following

observations.

First, global order conservation places a strong constraint on the classically nondimensionalized density profile in the bid-centered frame. We have seen that at $\epsilon \ll 1$, the density profile becomes concave upward near the bid, accounting for an increasing fraction of the allowed "remainder area" as $\epsilon \to 0$ (see Figs. 3 and 28). Since this remainder area is fixed at unity, it can be conserved only if the density profile approaches one more quickly with increasing price. Low density at low price appears to lead to more frequent persistent steps in the effective short-term random walk, and hence large short-term diffusivity. However, increased density far from the bid indicates less impact from market orders relative to the relaxation time of the Poisson distribution, and thus a lower long-time diffusivity.

The qualitative behavior of the bid-centered density profile is the same as that of the midpoint-centered profile, and this is expected because the spread distribution is stationary, rather than diffusive. In other words, the only way the diffusion of the bid or ask can differ from that of the midpoint is for the spread to either increase or decrease for several succeeding steps. Such autocorrelation of the spread cannot accumulate with time if the spread itself is to have a stationary distribution. Thus, the shift in the midpoint over some time interval can only differ from that of the bid or ask by at most a constant, as a result of a few correlated changes in the spread. This difference cannot grow with time, however, and so does not affect the diffusivity at long times.

Indeed, both of the predicted corrections to the classical estimate for diffusivity are seen in simulation results for midpoint diffusion. The simulation results, however, show that the implied autocorrelations change the diffusivity by factors of $\sqrt{\epsilon}$, suggesting that these corrections require a more subtle derivation than the one attempted here. This will be evidenced by the difficulty of obtaining a source term $\mathcal S$ in density coordinates (sectionIII F), which satisfied both the global order conservation law, and the proper zero-price boundary condition, in the midpoint-centered frame.

An interesting speculation is that the subtlety of these correlations also causes the density $n\left(p,t\right)$ in bid-centered coordinates not to approximate the mean-field condition at any of the parameters studied here, as noted in Sec. III D. Since short-term and long-term diffusivity corrections are related by a hard constraint, the difficulty of producing the late-time density profile should match that of producing the early-time profile. The midpoint-centered profile is potentially easier, in that the late-time complexity must be matched by a combination of the early-time density profile and the scaling of the expected spread. It appears that the complex scaling is absorbed in the spread, as per Fig. 10 and Fig. 24, leaving a density that can be approximately calculated with the methods used here.

F. Master equations and mean-field approximations

There are two natural limits in which functional configurations may become simple enough to be tractable probabilistically, with analytic methods. They correspond to mean field theories in which fluctuations of the dual differentials of either $N\left(p,t\right)$ or $p\left(N,t\right)$ are independent. In the first case, probabilities may be defined for any density $n\left(p,t\right)$ independently at each p, and in the second for the separation intervals $x\left(N,t\right)$ at each N. The mean field theory from the first approximation will be solved in Subsec. III F 1, and that from the second in Subsec. III G. As mentioned above, because the fluctuation independence approximation is only usable in a midpoint-centered frame, $n\left(p,t\right)$ will refer always to this frame. $x\left(N,t\right)$ is well-defined without reference to any frame.

1. A number density master equation

If share-number fluctuations are independent at different p, a density $\pi\left(n,p,t\right)$ may be defined, which gives the probability to find n orders in bin (p,p+dp), at time t. The normalization condition defining π as a probability density is

$$\sum_{n} \pi(n, p, t) = 1, \tag{13}$$

for each bin index p and at every t. The index t will be suppressed henceforth in the notation since we are looking for time-independent solutions.

Supposing an arbitrary density of order-book configurations $\pi(n,p)$ at time t, the stochastic dynamics of the configurations causes probability to be redistributed according to the master equation

$$\frac{\partial}{\partial t} \quad \pi(n,p) = \frac{\alpha(p) dp}{\sigma} \left[\pi(n-\sigma,p) - \pi(n,p) \right]
+ \frac{\delta}{\sigma} \left[(n+\sigma) \pi(n+\sigma,p) - n\pi(n,p) \right]
+ \frac{\mu(p)}{2\sigma} \left[\pi(n+\sigma,p) - \pi(n,p) \right]
+ \sum_{\Delta p} P_{+}(\Delta p) \left[\pi(n,p-\Delta p) - \pi(n,p) \right]
+ \sum_{\Delta p} P_{-}(\Delta p) \left[\pi(n,p+\Delta p) - \pi(n,p) \right]. (14)$$

Here $\partial \pi (n,p)/\partial t$ is a continuum notation for $[\pi (n,p,t+\delta t)-\pi (n,p,t)]/\delta t$, where δt is an elementary time step, chosen short enough that at most one event alters any typical configuration. Eq. (14) represents a general balance between additions and removals, without regard to the meaning of n. Thus, $\alpha (p)$

is a function that must be determined self-consistently with the choice of frame. As an example of how this works, in a bid-centered frame, $\alpha(p)$ takes a fixed value $\alpha(\infty)$ at all p, because the deposition rate is independent of position and frame shifts. The midpoint-centered frame is more complicated, because depositions below the midpoint cause shifts that leave the deposited order above the midpoint. The specific consequence for $\alpha(p)$ in this case will be considered below. $\mu(p)/2$ is. similarly, the rate of market orders surviving to cancel limit orders at price p. $\mu(p)/2$ decreases from $\mu(0)/2$ at the ask (for buy market orders, because μ total orders are divided evenly between buys and sells) to zero as $p \to \infty$, as market orders are screened probabilistically by intervening limit orders. $\alpha(\infty)$ and $\mu(0)$ are thus the parameters α and μ of the simulation.

The lines of Eq. (14) correspond to the following events. The term proportional to $\alpha(p) dp/\sigma$ describes depositions of discrete orders at that rate (because α is expressed in shares per price per time), which raise configurations from $n-\sigma$ to n shares at price p. The term proportional to δ comes from deletions and has the opposite effect, and is proportional to n/σ , the number of orders that can independently decay. The term proportional to $\mu(p)/2\sigma$ describes market order annihilations. For general configurations, the preceding three effects may lead to shifts of the origin by arbitrary intervals Δp , and P_{\pm} are for the moment unknown distributions over the frequency of those shifts. They must be determined self-consistently with the configuration of the book which emerges from any solution to Eq. (14).

A limitation of the simple product representation of frame shifts is that it assumes that whole order-book configurations are transported under $p \pm \Delta p \rightarrow p$, independently of the value of n(p). As long as fluctuations are independent, this is a good approximation for orders at all p which are not either the bid or the ask, either before or after the event that causes the shift. The correlations are never ignorable for the bins which are the bid and ask, though, and there is some distribution of instances in which any p of interest plays those parts. Approximate methods to incorporate those correlations will require replacing the product form with a sum of products conditioned on states of the order book, as will be derived below.

The important point is that the order-flow dependence of Eq. (14) is independent of these self-consistency requirements, and may be solved by use of generating functionals at general $\alpha\left(p\right)$, $\mu\left(p\right)$, and P_{\pm} . The solution, exact but not analytically tractable at general dp, will be derived in closed form in the next subsection. It has a well-behaved continuum limit at $dp \rightarrow 0$, however, which is analytically tractable, so that special case will be considered in the following subsection.

2. Solution by generating functional

The moment generating functional for π is defined for a parameter $\lambda \in [0, 1]$, as

$$\Pi(\lambda, p) \equiv \sum_{n/\sigma=0}^{\infty} \lambda^{n/\sigma} \pi(n, p).$$
 (15)

Introducing a shorthand for its value at $\lambda = 0$,

$$\Pi(0,p) = \pi(0,p) \equiv \pi_0(p), \qquad (16)$$

while the normalization condition (13) for probabilities gives

$$\Pi(1,p) = 1, \forall p. \tag{17}$$

By definition of the average of n(p) in the distribution π , denoted $\langle n(p) \rangle$,

$$\frac{\partial}{\partial \lambda} \Pi(\lambda, p) \bigg|_{\lambda \to 1} = \frac{\langle n(p) \rangle}{\sigma},$$
 (18)

and because Π will be regular in some sufficiently small neighborhood of $\lambda = 1$, one can expand

$$\Pi(\lambda, p) = 1 + (\lambda - 1) \frac{\langle n(p) \rangle}{\sigma} + \mathcal{O}(\lambda - 1)^{2}.$$
 (19)

Multiplying Eq. (14) by $\lambda^{n/\sigma}$ and summing over n, (and suppressing the argument p in the notation everywhere; $\alpha(p)$ or $\alpha(0)$ will be used where the distinction of the function from its boundary value is needed) the stationary solution for Π must satisfy

$$0 = \frac{\lambda - 1}{\sigma} \left\{ \alpha dp \Pi - \delta \sigma \frac{\partial \Pi}{\partial \lambda} - \frac{\mu}{2\lambda} (\Pi - \pi_0) \right\} \Big|_{(\lambda, p)}$$

$$+ \sum_{\Delta p} P_+ (\Delta p) \left[\Pi (\lambda, p - \Delta p) - \Pi (\lambda, p) \right]$$

$$+ \sum_{\Delta p} P_- (\Delta p) \left[\Pi (\lambda, p + \Delta p) - \Pi (\lambda, p) \right]. \tag{20}$$

Only the symmetric case with no net drift will be considered here for simplicity, which requires $P_{+}(\Delta p) = P_{-}(\Delta p) \equiv P(\Delta p)$. In a Fokker-Planck expansion, the (unrenormalized) diffusivity of whatever reference price is used as coordinate origin, is related to the distribution P by

$$D \equiv \sum_{\Delta p} P(\Delta p) \, \Delta p^2. \tag{21}$$

The rate at which shift events happen is

$$R \equiv \sum_{\Delta p} P(\Delta p) , \qquad (22)$$

and the mean shift amount appearing at linear order in derivatives (relevant at $p \to 0$), is

$$\langle \Delta p \rangle \equiv \frac{\sum_{\Delta p} P(\Delta p) \, \Delta p}{\sum_{\Delta p} P(\Delta p)}.$$
 (23)

Anywhere in the interior of the price range (where p is not at any stage the bid, ask, or a point in the spread), Eq. (20) may be written

$$\left\{ \frac{\partial}{\partial \lambda} - \frac{D}{\delta (\lambda - 1)} \frac{\partial^2}{\partial p^2} - \frac{\alpha dp - \mu/2\lambda}{\delta \sigma} \right\} \Pi = \frac{\mu}{2\delta \sigma \lambda} \pi_0.$$
(24)

Evaluated at $\lambda \to 1$, with the use of the expansion (19), this becomes

$$\left(1 - \frac{D}{\delta} \frac{\partial}{\partial p^2}\right) \langle n \rangle = \frac{\alpha dp}{\delta} - \frac{\mu}{2\delta} (1 - \pi_0).$$
(25)

At this point it is convenient to specialize to the case $dp \to 0$, wherein the eligible values of any $\langle n(p) \rangle$ become just σ and zero. The expectation is then related to the probability of zero occupancy (at each p) as

$$\langle n \rangle = \sigma \left[1 - \pi_0 \right], \tag{26}$$

yielding immediately

$$\frac{\alpha dp}{\delta} = \left[\frac{\mu}{2\delta \sigma} + \left(1 - \frac{D}{\delta} \frac{d^2}{dp^2} \right) \right] \langle n \rangle. \tag{27}$$

Eq. (27) defines the general solution $\langle n(p) \rangle$ for the master equation (14), in the continuum limit $2\alpha dp/\mu \rightarrow 0$. The shift distribution $P\left(\Delta p\right)$ appears only through the diffusivity D, which must be solved self-consistently, along with the otherwise arbitrary functions α and μ . The more general solution at large dp is carried out in App. B 1.

A first step toward nondimensionalization may be taken by writing Eq. (27) in the form (re-introducing the indexing of the functions)

$$\frac{\alpha(p)}{\alpha(\infty)} = \left[\frac{\mu(p)}{\mu(0)} + \epsilon \left(1 - \frac{D}{\delta} \frac{d^2}{dp^2}\right)\right] \frac{1}{\epsilon} \frac{\delta \langle n \rangle}{\alpha dp}.$$
 (28)

Far from the midpoint, where only depositions and cancellations take place, orders in bins of width dp are Poisson distributed with mean $\alpha (\infty) dp/\delta$. Thus, the asymptotic value of $\delta \langle n \rangle / \alpha (\infty) dp$ at large p is unity. This is consistent with a limit for $\alpha (p) / \alpha (\infty)$ of unity, and a limit for the screened $\mu (p) / \mu (0)$ of zero. The reason for grouping the nondimensionalized number density with $1/\epsilon$, together with the proper normalization of the characteristic price scale, will come from examining the decay of the dimensionless function $\mu (p) / \mu (0)$.

3. Screening of the market-order rate

In the context of independent fluctuations, Eq. (26) implies a relation between the mean density and the rate at which market orders are screened as price increases. The effect of a limit order, resident in the price bin p when a market order survives to reach that bin, is to prevent its arriving at the bin at p+dp. Though the nature of the shift induced, when such annihilation occurs, depends

on the comoving frame being modeled, the change in the number of orders surviving is independent of frame, and is given by

$$d\mu = -\mu \left(1 - \pi_0\right) = -\mu \left\langle n \right\rangle / \sigma. \tag{29}$$

Eq. (29) may be rewritten

$$\frac{d\log\left(\mu\left(p\right)/\mu\left(0\right)\right)=}{dp}-\frac{1}{\epsilon}\left(\frac{2\alpha\left(\infty\right)}{\mu\left(0\right)}\right)\left(\frac{\delta\left\langle n\left(p\right)\right\rangle}{\alpha\left(\infty\right)dp}\right),\ (30)$$

identifying the characteristic scale for prices as $p_c = \mu(0)/2\alpha(\infty) \equiv \mu/2\alpha$. Writing $\hat{p} \equiv p/p_c$, the function that screens market orders is the same as the argument of Eq. (28), and will be denoted

$$\frac{1}{\epsilon} \frac{\delta \langle n(p) \rangle}{\alpha(\infty) dp} \equiv \psi(\hat{p}) \tag{31}$$

Defining a nondimensionalized diffusivity $\beta \equiv D/\delta p_c^2$, Eq. (27) can then be put in the form

$$\frac{\alpha(p)}{\alpha(\infty)} = \left[\frac{\mu(p)}{\mu(0)} + \epsilon \left(1 - \beta \frac{d^2}{d\hat{p}^2}\right)\right] \psi, \tag{32}$$

with

$$\frac{\mu(p)}{\mu(0)} \equiv \varphi(\hat{p}) = \exp\left(-\int_0^{\hat{p}} d\hat{p}' \psi(\hat{p}')\right), \quad (33)$$

4. Verifying the conservation laws

Since nothing about the derivation so far has made explicit use of the frame in which n is averaged, the combination of Eq. (32) with Eq. (33) respects the conservation laws (9) and (10), if appropriate forms are chosen for the deposition rate $\alpha(p)$.

For example, in the bid-centered frame, $\alpha(p)/\alpha(\infty) = 1$ everywhere. Multiplying Eq. (32) by $d\hat{p}$ and integrating over the whole range from the bid to $+\infty$, we recover the nondimensionalized form of Eq. (9):

$$\int_0^\infty d\hat{p} \left(1 - \epsilon \psi\right) = 1,\tag{34}$$

iff we are careful with one convention. The integral of the diffusion term formally produces the first derivative $d\psi/d\hat{p}|_0^{\infty}$. We must regard this as a true first derivative, and consider its evaluation at zero continued far enough below the bid to capture the identically zero first derivative of the sell order depth profile.

In the midpoint centered frame, the correct form for the source term should be $\alpha(\hat{p})/\alpha(\hat{\infty}) = 1 + \Pr(\hat{s}/2 \geq \hat{p})$, whatever the expression for the cumulative distribution function. Recognizing that the integral of the CDF is, by parts, the mean value of $\hat{s}/2$, the same integration of Eq. (32) gives

$$\int_0^\infty d\hat{p} \left(1 - \epsilon \psi\right) = 1 - \frac{\langle \hat{s} \rangle}{2},\tag{35}$$

the nondimensionalized form of Eq. (10). Again, this works only if the surface contribution from integrating the diffusion term vanishes.

Neither of these results required the assumption of independent fluctuations, though that will be used below to give a simple approximate form for $\Pr(\hat{s}/2 \geq \hat{p}) \approx \varphi(\hat{p})$. They therefore provide a check that the extinction form (33) propagates market orders correctly into the interior of the order-book distribution, to respect global conservation. They also check the consistency of the intuitively plausible form for α in the midpoint-centered frame. The detailed form is then justified whenever the assumption of independent fluctuations is checked to be valid.

5. Self-consistent parametrization

The assumption of independent fluctuations of $n\left(p\right)$ used above to derive the screening of market orders, is equivalent to a specification of the CDF of the ask. Market orders are only removed between prices p and p+dp in those instances when the ask is at p. Therefore

$$\Pr\left(\hat{s}/2 \ge \hat{p}\right) = \varphi\left(\hat{p}\right),\tag{36}$$

the continuum limit of Eq. (12). Together with the form $\alpha(\hat{p})/\alpha(\hat{\infty}) = 1 + \Pr(\hat{s}/2 \ge \hat{p})$, Eq. (32) becomes

$$1 + \varphi = -\left[\frac{d\varphi}{d\hat{p}} + \epsilon \left(1 - \beta \frac{d^2}{d\hat{p}^2}\right) \frac{d\log\varphi}{d\hat{p}}\right]. \tag{37}$$

(If the assumption of independent fluctuations were valid in the bid-centered frame, it would take the same form, but with φ removed on the left-hand side.)

To consistently use the diffusion approximation, with the realization that for p=0, $n\pi\left(n,p-\Delta p\right)=0$ for essentially all Δp in Eq. (14), it is necessary to set the Fokker-Planck approximation to $\psi\left(0-\langle\Delta p\rangle\right)=0$ as a boundary condition. Nondimensionalized, this gives

$$\frac{\beta}{2} \frac{d^2 \psi}{d\hat{p}^2} \bigg|_{0} = \frac{R}{\delta} \left(\langle \Delta \hat{p} \rangle \frac{d}{d\hat{p}} - 1 \right) \psi \bigg|_{0}, \tag{38}$$

where R is the rate at which shifts occur (Eq. 22). In the solutions below, the curvature will typically be much smaller than ψ (0) \sim 1, so it will be convenient to enforce the simpler condition

$$\langle \Delta \hat{p} \rangle \left. \frac{d\psi}{d\hat{p}} \right|_{0} - \psi \left(0 \right) \approx 0,$$
 (39)

and verify that it is consistent once solutions have been evaluated.

Self-consistent expressions for β and $\langle \Delta p \rangle$ are then constructed as follows. Given an ask at some position a (in the midpoint-centered frame), there is a range from -a to a in which sell limit orders may be placed, which will induce positive midpoint-shifts. The shift amount is half

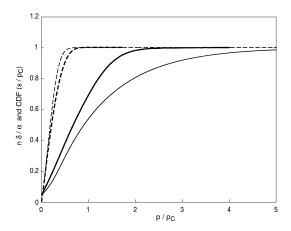


FIG. 19: Fit of the self-consistent solution with diffusivity term to simulation results for the midpoint-centered frame. Thin solid line is the analytic solution for the mean number density, and thick solid line is simulation result, at $\epsilon = 0.02$. Thin dashed line is the analytic prediction for the cumulative distribution function $\Pr(\hat{s}/2 \leq \hat{p})$, and thick dashed line is simulation result.

as great as the distance from the bid, so the measure for shifts $dP_+(\Delta p)$ from sell limit-order addition inherits a term $2\alpha\left(0\right)\left(d\Delta p\right)\Pr\left(a\geq\Delta p\right)$, where the last factor counts the instances with asks large enough to admit shifts by Δp . There is an equal contribution to dP_- from addition of buy limit orders. Symmetry requires that for every positive shift due to an addition, there is a negative shift due to evaporation with equal measure, so the contribution from buy limit order removal should equal that for sell limit order addition. When these contributions are summed, the measures for positive and negative shifts both equal

$$dP_{\pm}(\Delta p) = 4\alpha(\infty) (d\Delta p) \Pr(a \ge \Delta p). \tag{40}$$

Eq. (40) may be inserted into the continuum limit of the definition (21) for D, and then nondimensionalized to give

$$\beta = \frac{4}{\epsilon} \int_0^\infty d\Delta \hat{p} \left(\Delta \hat{p}\right)^2 \varphi \left(\Delta \hat{p}\right), \tag{41}$$

where the mean-field substitution of $\varphi(\Delta \hat{p})$ for $\Pr(a \ge \Delta p)$ has been used. Similarly, the mean shift amount used in Eq. (39) is

$$\langle \Delta \hat{p} \rangle = \frac{\int_0^\infty d\Delta \hat{p} \, (\Delta \hat{p}) \, \varphi \, (\Delta \hat{p})}{\int_0^\infty d\Delta \hat{p} \, \varphi \, (\Delta \hat{p}) \,.} \tag{42}$$

A fit of Eq. (37) to simulations, using these self-consistent measures for shifts, is shown in Fig. 19. This solution is actually a compromise between approximations with opposing ranges of validity. The diffusion equation using the mean order depth describes nonzero transport of limit orders through the midpoint, an approximation inconsistent with the correlations of shifts

with states of the order book. This approximation is a small error only at $\epsilon \to 0$. On the other hand, both the form of α , and the self-consistent solutions for $\langle \Delta \hat{p} \rangle$ and β , made use of the mean-field approximation, which we saw was only valid for $\epsilon \lesssim 1$. The two approximations appear to create roughly compensating errors in the intermediate range $\epsilon \sim 0.02$.

6. Accounting for correlations

The numerical integral implementing the diffusion solution actually doesn't satisfy the global conservation condition that the diffusion term integrate to zero over the whole price range. Thus, it describes diffusive transport of orders through the midpoint, and as such also doesn't have the right $\hat{p}=0$ boundary condition. The effective absorbing boundary represented by the pure diffusion solution corresponds roughly to the approximation made by Bouchaud $et\ al.\ [23]$ It differs from theirs, though, in that their method of images effectively approximates the region of the spread as a point, whereas Eq. (32) actually resolves the screening of market orders as the spread fluctuates.

Treating the spread region – roughly defined as the range over which market orders are screened – as a point is consistent with treating the resulting coarse-grained "midpoint" as an absorbing boundary. If the spread is resolved, however, it is not consistent for diffusion to transport any finite number density through the midpoint, because the midpoint is strictly always in the center of an open set with no orders, in a continuous price space. The correct behavior in a neighborhood of the "fine-grained midpoint" can be obtained by explicitly accounting for the correlation of the state of orders, with the shifts that are produced when market or limit order additions occur.

We expect the problem of recovering both the global conservation law and the correct $\hat{p}=0$ boundary condition to be difficult, as it should be responsible for the non-trivial corrections to short-term and long-term diffusion mentioned earlier. We have found, however, that by explicitly sacrificing the global conservation law, we can incorporate the dependence of shifts on the position of the ask, in an interesting range around the midpoint. At general ϵ , the corrections to diffusion reproduce the mean density over the main support of the CDF of the spread. While the resulting density does not predict that CDF (due to correlated fluctuations), it closely enough resembles the real density that the independent CDFs of the two are similar.

7. Generalizing the shift-induced source terms

Nondimensionalizing the generating-functional master equation (20) and keeping leading terms in dp at $\lambda \to 1$,

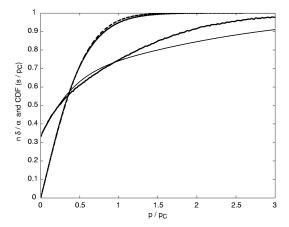


FIG. 20: Reconstruction with source terms \mathcal{S} that approximately account for correlated fluctuations near the midpoint. $\epsilon=0.2$. Thick solid line is averaged order book depth from simulations, and thin solid is the mean field result. Thin dotted line is the simulated CDF for $\hat{s}/2$, and thick dotted line is the mean field result. Thick dashed line is the CDF that would be produced from the simulated depth, if the mean-field approximation were exact.

get

$$\frac{\alpha(\hat{p})}{\alpha(\hat{\infty})} = \left(\frac{\mu(\hat{p})}{\mu(\hat{0})} + \epsilon\right) \psi(\hat{p})$$

$$- \int dP_{+} (\Delta \hat{p}) \left[\psi(\hat{p} - \Delta \hat{p}) - \psi(\hat{p})\right]$$

$$- \int dP_{-} (\Delta \hat{p}) \left[\psi(\hat{p} + \Delta \hat{p}) - \psi(\hat{p})\right] (43)$$

where $dP_{\pm}(\Delta \hat{p})$ is the nondimensionalized measure that results from taking the continuum limit of P_{\pm} in the variable $\Delta \hat{p}$.

Eq. (43) is inaccurate because the number of orders shifted into or out of a price bin p, at a given spread, may be identically zero, rather than the unconditional mean value ψ . We take that into account by replacing the last two lines of Eq. (43) with lists of source terms, whose forms depend on the position of the ask, weighted by the probability density for that ask. Independent fluctuations are assumed by using Eq. (36).

It is convenient at this point to denote the replacement of the last two lines of Eq. (43) with the notation S, yielding

$$\frac{\alpha\left(\hat{p}\right)}{\alpha\left(\hat{\infty}\right)} = \left(\frac{\mu\left(\hat{p}\right)}{\mu\left(\hat{0}\right)} + \epsilon\right)\psi - \mathcal{S}.\tag{44}$$

The global conservation laws for orders would be satisfied if $\int d\hat{p} S = 0$.

The source term S is derived approximately in App. B2. The solution to Eq. (44) at $\epsilon = 0.2$, with the simple-diffusive source term replaced by the evaluations (B29 - B37), is compared to the simulated order-

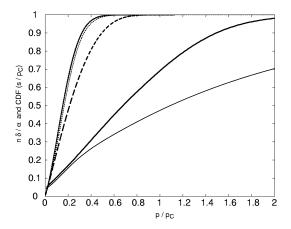


FIG. 21: Reconstruction with correlated source terms for $\epsilon = 0.02$. Line style and thickness are the same as in Fig. 20.

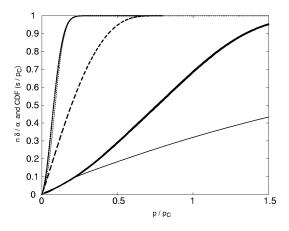


FIG. 22: Reconstruction with correlated source terms for $\epsilon=0.002$. Line style and thickness are the same as in Fig. 20.

book depth and spread distribution in Fig. 20. The simulated $\langle n\left(p\right)\rangle$ satisfies Eq. (35), showing what is the correct "remainder area" below the line $\langle n\rangle\equiv 1$. The numerical integral deviates from that value by the incorrect integral $\int d\hat{p}\mathcal{S}\neq 0$. However, most of the probability for the spread lies within the range where the source terms \mathcal{S} are approximately correct, and as a result the distribution for $\hat{s}/2$ is predicted fairly well.

Even where the mean-field approximation is known to be inadequate, the source terms defined here capture most of the behavior of the order-book distribution in the region that affects the spread distribution. Fig. 21 shows the comparison to simulations for $\epsilon=0.02$, and Fig. 22 for $\epsilon=0.002$. Both cases fail to reproduce the distribution for the spread, and also fail to capture the large- \hat{p} behavior of ψ . However, they approximate ψ at small \hat{p} well enough that the resulting distribution for the spread is close to what would be produced by the simulated ψ if fluctuations were independent.

G. A mean-field theory of order separation intervals: The Independent Interval Approximation

A simplifying assumption that is in some sense dual to independent fluctuations of n(p), is independent fluctuations in the intervals x(N) at different N. Here we develop a mean-field theory for the order separation intervals in this model. From this, we will also be able to make an estimate of the depth profiles for any value of the parameters. For convenience of notation we will use x_N to denote x(N).

Limit order placements are considered to take place strictly on sites which are not occupied. This is the same level of approximation as made in the previous section. The time step is normalized to unity, as above, so that rates are equal to probabilities after one update of the whole configuration. The rates α and μ used in this section correspond to $\alpha(\infty)$ and $\mu(0)$ as defined earlier.

As shown in Fig. 15 the configuration is entirely specified instant by instant if the instantaneous values of the order separation intervals are known.

Consider now, how these intervals might change due to various processes. For the spread x_0 , these processes and the corresponding change in x_0 , are listed below.

- 1. $x_0 \to x_0 + x_1$ with rate $(\delta + \mu/2\sigma)$ (when the ask either evaporates or is deleted by a market order).
- 2. $x_0 \to x_0 + x_{-1}$ with rate $(\delta + \mu/2\sigma)$ (when the bid either evaporates or is deleted by the corresponding market order).
- 3. $x_0 \to x'$ for any value $1 \le x' \le x_0 1$, when a sell limit order is deposited anywhere in the spread. The rate for any single deposition is $\alpha dp/\sigma$, so the cumulative rate for *some* deposition is $\alpha dp(x_0 1)/\sigma$. (The -1 comes from the prohibition against depositing on occupied sites.)
- 4. Similarly $x_0 \to x_0 x'$ for any $1 \le x' \le x_0 1$, when a buy limit order is deposited in the spread, also with cumulative rate $\alpha dp(x_0 1)/\sigma$.
- 5. Since the above processes describe all possible single-event changes to the configuration, the probability that it remains unchanged in a single time step is $1 2\delta \mu/\sigma 2\alpha dp (x_0 1)/\sigma$.

In all that follows, we will put $\sigma = 1$ without loss of generality. If we know x_0 , x_1 , and x_{-1} at time t, the expected value at time t + dt is then

$$x_{0}(t+dt) = x_{0}(t) \left[1 - 2\delta - \mu_{0} - 2\alpha (x_{0}-1)\right] + (x_{0} + x_{1}) \left(\delta + \frac{\mu}{2}\right) + (x_{0} + x_{-1}) \left(\delta + \frac{\mu}{2}\right) + (\alpha_{0}dp)x_{0}(x_{0}-1)$$

$$(45)$$

Here, $x_i(t)$ represents the value of the interval averaged over many realizations of the process evolved up to time t.

Again representing the finite difference as a time derivative, the change in the expected value, given x_0 , x_1 , and x_{-1} , is

$$\frac{dx_0}{dt} = (x_1 + x_{-1}) \left(\delta + \frac{\mu}{2} \right) - (\alpha dp) x_0 \left(x_0 - 1 \right). \tag{46}$$

Were it not for the quadratic term arising from deposition, Eq. (46) would be a linear function of x_0 , x_1 , and x_{-1} . However we now need an approximation for $\langle x_0^2 \rangle$, where the angle brackets represent an average over realizations as before or equivalently a time average in the steady state. Let us for the moment assume that we can approximate $\langle x_0^2 \rangle$ by $a \langle x_0 \rangle^2$, where a is some as yet undetermined constant to be determined self-consistently. We will make this approximation for all the x_k 's. This is clearly not entirely accurate because the PDF of x_k could depend on k (as indeed it does. We will comment on this a little later). However as we will see this is still a very good approximation.

We will therefore make this approximation in Eq. (46) and everywhere below, and look for steady state solutions when the x_k 's have reached a time independent average value.

It then follows that,

$$(\delta + \mu/2)(x_1 + x_{-1}) = a\alpha dp x_0(x_0 - 1)$$
 (47)

The interval x_k may be though of as the inverse of the density at a distance $\sum_{j=0}^{k-1} x_j$ from the bid. That is, $x_i \approx 1/\left\langle n\left(\sum_{j=0}^{i-1} x_j dp\right)\right\rangle$, the dual to the mean depth, at least at large i. It therefore makes sense to introduce a normalized interval

$$\hat{x}_i \equiv \epsilon \frac{\alpha}{\delta} x_i dp = \frac{x_i dp}{p_c} \approx \frac{1}{\psi\left(\sum_{j=0}^{i-1} \hat{x}_j\right)},\tag{48}$$

the mean-field inverse of the normalized depth ψ . In this nondimensionalized form, Eq. (47) becomes

$$(1+\epsilon)(\hat{x}_1 + \hat{x}_{-1}) = a\hat{x}_0(\hat{x}_0 - d\hat{p}), \tag{49}$$

where $d\hat{p} = dp/p_c$.

Since the depth profile is symmetric about the origin, $\hat{x}_1 = \hat{x}_{-1}$. From the equations, it can be seen that this ansatz is self-consistent and extends to all higher \hat{x}_i . Substituting this in Eq. (49) we get

$$(1+\epsilon)\,\hat{x}_1 = \frac{a}{2}\hat{x}_0\,(\hat{x}_0 - d\hat{p}) = (1+\epsilon)\,\hat{x}_{-1} \qquad (50)$$

Proceeding to the change of x_1 , the events that can occur, with their probabilities, are shown in Table V, with the remaining probability that x_1 remains unchanged.

The differential equation for the mean change of x_1 can be derived along previous lines and becomes

$$\frac{dx_1}{dt} = \left(2\delta + \frac{\mu}{2}\right)x_2 - \left(\delta + \frac{\mu}{2}\right)x_1 + \alpha dp \left[\frac{x_0(x_0 - 1)}{2} - \frac{x_1(x_1 - 1)}{2} - x_1(x_0 - 1)\right]$$
(51)

case	rate	range
$x_1 \to x_2$	$(\delta + \mu/2)$	
$x_1 \to (x_1 + x_2)$	δ	
$x_1 \to x'$	αdp	$x' \in (1, x_0 - 1)$
$x_1 \to x_1 - x'$	αdp	$x' \in (1, x_1 - 1)$

TABLE V: Events that can change the value of x_1 , with their rates of occurrence.

Note that in the above equations, the mean-field approximation consists of assuming that terms like $\langle x_0 x_1 \rangle$ are approximated by the product $\langle x_0 \rangle \langle x_1 \rangle$. This is thus an 'independent interval' approximation.

Nondimensionalizing Eq. (51) and combining the result with Eq. (50) gives the stationary value for x_2 from x_0 and x_1 ,

$$(1+2\epsilon)\,\hat{x}_2 = \frac{a}{2}\hat{x}_1\,(\hat{x}_1 - d\hat{p}) + \hat{x}_1\,(\hat{x}_0 - d\hat{p})\,. \tag{52}$$

Following the same procedure for general k, the nondimensionalized recursion relation is

$$(1 + k\epsilon) \hat{x}_k = \frac{a}{2} \hat{x}_{k-1} (\hat{x}_{k-1} - d\hat{p}) + \hat{x}_{k-1} \sum_{i=0}^{k-2} (\hat{x}_i - d\hat{p}).$$
(53)

1. Asymptotes and conservation rules

Far from the bid or ask, \hat{x}_k must go to a constant value, which we denote \hat{x}_{∞} . In other words, for large k, $\hat{x}_{k+1} \to \hat{x}_k$. Taking the difference of Equation (53) for k+1 and k in this limit gives the identification

$$\epsilon \hat{x}_{\infty} = \hat{x}_{\infty} \left(\hat{x}_{\infty} - d\hat{p} \right), \tag{54}$$

or $\hat{x}_{\infty} = \epsilon + d\hat{p}$. Apart from the factor of $d\hat{p}$, arising from the exclusion of deposition on already-occupied sites, this agrees with the limit $\psi\left(\infty\right) \to 1/\epsilon$ found earlier. In the continuum limit $d\hat{p} \to 0$ at fixed ϵ , these are the same.

From the large-k limit of Eq. (53), one can also solve easily for the quantity $S_{\infty} \equiv \sum_{i=0}^{\infty} (\hat{x}_i - \hat{x}_{\infty})$, which is related to the bid-centered order conservation law mentioned in Section III C. Dividing by a factor of \hat{x}_{∞} at large k,

$$(1 + k\epsilon) = \frac{a}{2} (\hat{x}_{\infty} - d\hat{p}) + \sum_{i=0}^{k-2} (\hat{x}_i - d\hat{p}), \qquad (55)$$

or, using Eq. (54) and rewriting the sum on the righthand side as $\sum_{i=0}^{k-2} (\hat{x}_i - \hat{x}_\infty) + \sum_{i=0}^{k-2} \hat{x}_\infty - d\hat{p}$

$$1 + (1 - \frac{a}{2})\epsilon = S_{\infty}. (56)$$

The interpretation of S_{∞} is straightforward. There are k+1 orders in the price range $\sum_{i=0}^{k} x_i$. Their decay rate

ϵ	S_{∞} from theory	S_{∞} from MCS
0.66	1	1.000
0.2	1	1.000
0.04	1	0.998
0.66 0.2 0.04 0.02	1	1.000

TABLE VI: Theoretical vs. results from simulations for S_{∞} .

is $\delta(k+1)$, and the rate of annihilation from market orders is $\mu/2$. The rate of additions, up to an uncertainty about what should be considered the center of the interval, is $(\alpha dp) \sum_{i=0}^{k} (x_i - 1)$ in the bid-centered frame (where effective α is constant and additions on top of previously occupied sites is forbidden). Equality of addition and removal is the bid-centered order conservation law (again), in the form

$$\frac{\mu}{2} + \delta(k+1) = \alpha dp \sum_{i=0}^{k} (x_i - 1).$$
 (57)

Taking k large, nondimensionalizing, and using Eq. (54), Eq. (57) becomes

$$1 = S_{\infty}.\tag{58}$$

This conservation law is indeed respected to a remarkable accuracy in Monte Carlo simulations of the model as indicated in table VI.

In order that the equation for the x's obey this exact conservation law, we require Eq. 56 to be equal to Eq. 58. We can hence now self-consistently set the value of a=2

The value of a implies that we have now set $\langle x_k^2 \rangle \sim 2 \langle x_k \rangle^2$. This would be strictly true if the probability distribution function of the interval x_k were exponentially distributed for all k. This is generally a good approximation for large k for any ϵ . Fig 23 shows the numerical results from Monte Carlo simulations of the model, for the probability distribution function for three intervals x_0 , x_1 and x_5 at $\epsilon = 0.1$. The functional form for $P(x_0)$ and $P(x_1)$ are better approximated by a Gaussian than an exponential. However $P(x_5)$ is clearly an exponential.

Eq. (58) has an important consequence for the shortterm and long-term diffusivities, which can also be seen in simulations, as mentioned in earlier sections. The nondimensionalization of the diffusivity D with the rate parameters, suggests a classical scaling of the diffusivity

$$D \sim p_c^2 \delta = \frac{\mu^2}{4\alpha^2} \delta. \tag{59}$$

As mentioned earlier, it is observed from simulations that the locally best short-time fit to the actual diffusivity of the midpoint is $\sim \sqrt{1/\epsilon}$ times the estimate (59), and the long-time diffusivity is $\sim \sqrt{\epsilon}$ times the classical estimate. While we do not yet know how to derive this relation analytically, the fact that early and late-time renormalizations must have this qualitative relation can be argued from the conservation law (58).

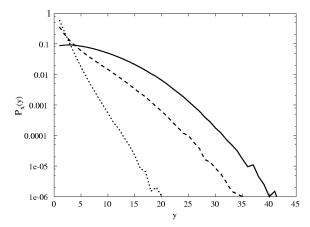


FIG. 23: The probability distribution functions $P_x(y)$ vs. y for the intervals $x = x_0, x_1$ and x_5 at $\epsilon = 0.1$, on a semi-log scale. Solid curve is for x_0 , dashed for x_1 , and dot for x_5 . The functional form of the distribution changes from a Gaussian to an exponential.

 S_{∞} is the area enclosed between the actual density and the asymptotic value. Increases in $1/\epsilon$ (descaled marketorder rate) deplete orders near the spread, diminishing the mean depth at small \hat{p} , and induce the upward curvature seen in Fig. 3, and even more strongly in Fig. 28 below. As noted above, they cause more frequent shifts (more than compensating for the slight decrease in average step size), and increase the classically descaled diffusivity β . However, as a result, this increases the fraction of the area in S_{∞} accumulated near the spread, requiring that the mean depth at larger \hat{p} increase to compensate (see Fig. 28). The resulting steeper approach to the asymptotic depth at prices greater than the mean spread, and the larger negative curvature of the distribution, are fit by an effective diffusivity that decreases with increasing $1/\epsilon$. Since the distribution further from the midpoint represents the imprint of market order activity further in the past, this effective diffusivity describes the long-term evolution of the distribution. The resulting anticorrelation of the small- \hat{p} and large- \hat{p} effective diffusion constants implied by conservation of the area S_{∞} is exactly consistent with their respective $\sim \sqrt{1/\epsilon}$ and $\sim \sqrt{\epsilon}$ scalings. The general idea here is to connect diffusivities at short and long time scales to the depth profile near the spread and far away from the spread respectively. The conservation law for the depth profile, then implies a connection between these two diffusivities.

2. Direct simulation in interval coordinates

The set of equations determined by the general form (53) is ultimately parametrized by the single input \hat{x}_0 . The correct value for \hat{x}_0 is determined when the \hat{x}_k are solved recursively, by requiring convergence to \hat{x}_{∞} . We do this recursion numerically, in the same manner as was done to solve the differential equation for

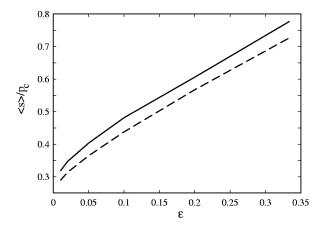


FIG. 24: The mean value of the spread in nondimensional units $\hat{s}=s/p_c$ as a function of ϵ . The numerical value above (solid) is compared with the theoretical estimate below (dash).

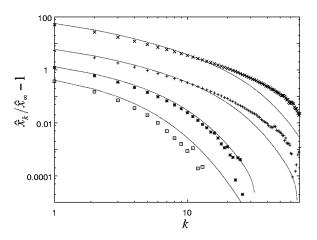


FIG. 25: Four pairs of curves for the quantity $\hat{x}_k/\hat{x}_\infty - 1$ vs. k. The value of ϵ increases from top to bottom ($\epsilon = 0.02, 0.04, 0.2, 0.66$). In each pair of curves, the markers are obtained from simulations while the solid curve is the prediction of Eq. 53 evaluated numerically. The difference between numerics and mean-field increases as ϵ decreases, especially for large k.

the normalized mean density $\psi(\hat{p})$.

In Fig. 24 we compare the numerical result for \hat{x}_0 with the analytical estimate generated as explained above. The results are surprisingly good throughout the entire range. Though the theoretical value consistently underestimates the numerical value, yet the functional form is captured accurately.

In Fig. 25, the values of x_k for all k, are compared to the values determined directly from simulations.

Fig. 26 shows the same data on a semilog scale for $\hat{x}_k/\hat{x}_{\infty} - 1$, showing the exponential decay at large argument characteristic of a simple diffusion solution. The IIA is clearly a good approximation for large ϵ . However for small ϵ it starts deviating significantly from the

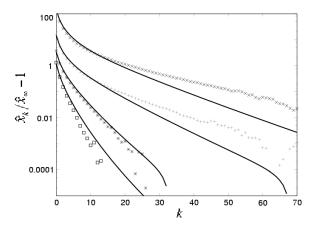


FIG. 26: Same plot as Fig. 25 but on a semi-log scale to show exponential decay at large k.

simulations, especially for large k.

The values of x_k computed from the IIA, can be very directly used to get an estimation of the price impact. The price impact, as defined in earlier sections, can be thought of as the change in the position of the midpoint (or the bid), consecutive to a certain number of orders being filled. Within the framework of the simplified model we study here, this is simply the quantity $\langle \Delta m \rangle = 1/2 \sum_{k'=1}^k x_{k'}$, for k orders. The factor of 1/2 comes from considering the change in the position of the midpoint and not the bid. Fig. 27 shows $\langle \Delta m \rangle$ nondimensionalized by p_c plotted as a function of the number of orders (multiplied by ϵ), for three different values of ϵ . Again, the theory matches quite well with the numerics, qualitatively. For large ϵ the agreement is quantitative as well.

The simplest approximation to the density profiles in the midpoint-centered frame is to continue to approximate the mean density as $1/x_k$, but to regard that density as evaluated at position $x_0/2+\sum_{k=1}^i x_k$. This clearly is not an adequate treatment in the range of the spread, both because the intervals are discrete, whereas mean ψ is continuous, and because the density profiles satisfy different global conservation laws associated with nonconstancy of α . For large k however, this approximation might hold. The mean-field values (only) corresponding to a plot of $\epsilon\psi(\hat{p})$ versus \hat{p} , are shown in Fig. 28. Here the theoretically estimated x_k 's at different parameter values are used to generate the depth profile using the procedure detailed above.

A comparison of the theoretically estimated profiles with the results from Monte Carlo simulations of the model, is shown in Fig. 29. As evident, the theoretical estimate for the density profile is better for large ϵ rather than small ϵ .

We can also generalize the above analysis to when the order placement process is no longer uniform. In particular it has been found that a power-law order placement process is relevant [23, 26]. We carry out the above analy-

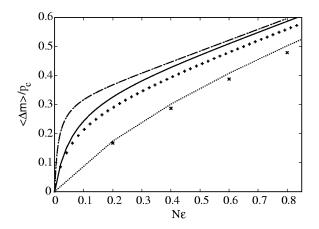


FIG. 27: Three pairs of curves for the quantity $\langle \Delta m \rangle/p_c$ vs. $N\epsilon$ where $\langle \Delta m \rangle = 1/2 \sum_{k=1}^N x_k$. The value of ϵ increases from top to bottom ($\epsilon = 0.002, 0.02, 0.2$). In each pair of curves, the markers are obtained from simulations while the solid curve is the prediction of the IIA. For $\epsilon = 0.002$, we show only the theoretical prediction. The theory captures the functional form of the price impact curves for different ϵ . Quantitatively, its better for larger epsilon, as remarked earlier.

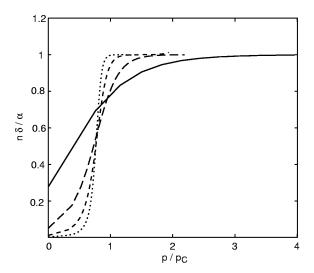


FIG. 28: Density profiles for different values of ϵ ranging over the values 0.2, 0.02, 0.004, 0.001, obtained from the Independent Interval Approximation.

sis for when $\alpha = {\Delta_0}^\beta/(\Delta + {\Delta_0})^\beta$ where Δ is the distance from the current bid and Δ_0 determines the 'shoulder' of the power-law. We find an interesting dependence of the existence of solutions on β . In particular we find that for $\beta > 1$, Δ_0 needs to be larger than some value (which depends on β as well as other parameters of the model such as μ and δ) for solutions of the IIA to exist. This might be interpreted as a market order wiping out the entire book, if the exponent is too large. When solutions exist, we find that the depth profile has a peak , consistent with the findings of [23]. In Fig. 30 the depth profiles

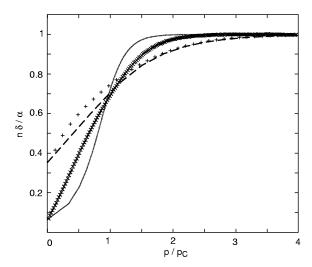


FIG. 29: Density profiles from Monte Carlo simulation (markers) and the Independent Interval Approximation (lines). Pluses and dash line are for $\epsilon=0.2$, while crosses and dotted line are for $\epsilon=0.02$.

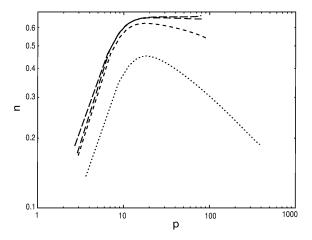


FIG. 30: Density profiles for a power-law order placement process for different values of Δ_0 .

for three different values of Δ_0 are plotted.

IV. CONCLUDING REMARKS

A. Ongoing work on empirical validation

Members of our group are working on the problem of empirically testing this model. We are using a dataset from the London Stock Exchange. We have chosen this data because it contains every order and every cancellation. This makes it possible to measure all the parameters of this model directly. It is also possible to reconstruct the order book and measure all the statistical properties we have studied in this paper. Our empirical work so far shows that, despite its many limitations, our model can act as an effective guide to future research. We believe that the main discrepancies between the predictions of our model and the data can be dealt with by using a more sophisticated model of order flow. We summarize some of the planned improvements in the following subsection.

B. Future Enhancements

As we have mentioned above, the zero intelligence, IID order flow model should be regarded as just a starting point from which to add more complex behaviors. We are considering several enhancements to the order flow process whose effects we intend to discuss in future papers. Some of the enhancements include:

• Trending of order flow.

We have demonstrated that IID order flow necessarily leads to non-IID prices. The converse is also true: Non-IID order flow is necessary for IID prices. In particular, the order flow must contain trends, i.e. if order flow has recently been skewed toward buying, it is more likely to continue to be skewed toward buying. If we assume perfect market efficiency, in the sense that prices are a random walk, this implies that there must be trends in order flow.

• Power law placement of limit prices

For both the London Stock Exchange and the Paris Bourse, the distribution of the limit price relative to the best bid or ask appears to decay as a power-law [23, 26]. Our investigations of this show that this can have an important effect. Exponents larger than one result in order books with a finite number of orders. In this case, depending on other parameters, there is a finite probability that a single market order can clear the entire book (see Section III G 2).

• Power law or log-normal order size distribution.

Real order placement processes have order size distributions that appear to be roughly like a lognormal distribution with a power law tail [27]. This has important effects on the fluctuations in liquidity.

• Non-Poisson order cancellation process.

When considered in real time order placement cancellation does not appear to be Poisson [21]. However, this may not be a bad approximation in event time rather than real time.

• Conditional order placement.

Agents may conditionally place larger market orders when the book is deeper, causing the market impact function to grow more slowly. We intend to measure this effect and incorporate it into our model.

• Feedback between order flow and prices.

In reality there are feedbacks between order flow and price movements, beyond the feedback in the reference point for limit order placement built into this model. This can induce bursts of trading, causing order flow rates to speed up or slow down, and give rise to clustered volatility.

The last item is just one of many examples of how one can surely improve the model by making order flow conditional on available information. However, we believe it is important to first gain an understanding of the properties of simple unconditional models, and then build on this foundation to get a fuller understanding of the problem.

C. Comparison to standard models based on valuation and information arrival

In the spirit of Gode and Sunder [3], we assume a simple, zero-intelligence model of agent behavior and show that the market institution exerts considerable power in shaping the properties of prices. While not disputing that agent behavior might be important, our model suggests that, at least on the short timescale many of the properties of the market are dictated by the market institution, and in particular the need to store supply and demand. Our model is stochastic and fully dynamic, and makes predictions that go beyond the realm of experimental economics, giving quantitative predictions about fundamental properties of a real market. We have developed what were previously conceptual toy models in the physics literature into a model with testable explanatory power.

This raises questions about the comparison to standard models based on the response of valuations to news. The idea that news might drive changes in order flow rates is compatible with our model. That is, news can drive changes in order flow, which in turn cause the best bid or ask price to change. But notice that in our model there are no assumptions about valuations. Instead, everything depends on order flow rates. For example, the diffusion rate of prices increases as the 5/2 power of market order flow rate, and thus volatility, which depends on the square root of the diffusion rate, increases as the 5/4power. Of course, order flow rates can respond to information; an increase in market order rate indicates added impatience, which might be driven by changes in valuation. But changes in long-term valuation could equally well cause an increase in limit order flow rate, which decreases volatility. Valuation per se does not determine whether volatility will increase or decrease. Our model says that volatility does not depend directly on valuations, but rather on the urgency with which they are felt, and the need for immediacy in responding to them.

Understanding the shape of the price impact function was one of the motivations that originally set this project into motion. The price impact function is closely related to supply and demand functions, which have been central aspects of economic theory since the 19th century. Our model suggests that the shape of price impact functions in modern markets is significantly influenced not so much by strategic thinking as by an economic fundamental: The need to store supply and demand in order to provide liquidity. A priori it is surprising that this requirement alone may be sufficient to dictate at least the broad outlines of the price impact curve.

Our model offers a "divide and conquer" strategy to understanding fundamental problems in economics. Rather than trying to ground our approach directly on assumptions of utility, we break the problem into two parts. We provide an understanding of how the statistical properties of prices respond to order flow rates, and leave the problem open of how order flow rates depend on more fundamental assumptions about information and utility. Order flow rates have the significant advantage that, unlike information, utility, or the cognitive powers of an agent, they are directly measurable. We hope that by breaking the problem into two pieces, and partially solving the second piece, we can ultimately help provide a deeper understanding of how markets work.

APPENDIX A: RELATIONSHIP OF PRICE IMPACT TO CUMULATIVE DEPTH

An important aspect of markets is the immediate liquidity, by which we mean the immediate response of prices to incoming market orders. When a market order enters, its execution range depends both on the spread and on the depth of the orders in the book. These determine the sequence of transaction prices produced by that order, as well as the instantaneous market impact. Long term liquidity depends on the longer term response of the limit order book, and is characterized by the price impact function $\phi(\omega, \tau)$ for values of $\tau > 0$. Immediate liquidity affects short term volatility, and long term liquidity affects volatility measured over longer timescales. In this section we address only short term liquidity. We address volatility on longer timescales in section II B 4.

We characterize liquidity in terms of either the depth profile or the price impact. The depth profile n(p,t) is the number of shares n at price p at time t. For many purposes it is convenient to think in terms of the cumulative depth profile N, which is the sum of n values up (or down) to some price. For convenience we establish a reference point at the center of the book where we define $p \equiv 0$ and $N(0) \equiv 0$). The reference point can be either the midpoint quote, or the best bid or ask. We also study the price impact function $\Delta p = \phi(\omega, \tau, t)$, where Δp is the shift in price at time $t + \tau$ caused by an order of size ω placed at time t. Typically we define Δp as the shift in the midpoint price, though it is also possible to use the best bid or ask (Eq. 1).

The price impact function and the depth profile are

closely related, but the relationship is not trivial. $N(\Delta p)$ gives us the average total number of orders upto a distance Δp away from the origin. Whereas, in order to calculate the price impact, what we need is the average shift Δp caused by a fixed number of orders. Making the identifications $p=\Delta p,\ N=\omega,$ and choosing a common reference point, the instantaneous price impact is the inverse of the instantaneous cumulative depth, i.e. $\phi(\omega,0,t)=N^{-1}(\omega,t).$ This relationship is clearly true instant by instant. However it is not true for averages, since the mean of the inverse is not in general equal to the inverse of the means, i.e. $\langle \phi \rangle \neq \langle N \rangle^{-1}.$ This is highly relevant here, since because the fluctuations in these functions are huge, our interest is primarily in their statistical properties, and in particular the first few moments.

A relationship between the moments can be derived as follows:

1. Moment expansion

There is some subtlety in how we relate the market impact to the cumulative order count $N\left(p,t\right)$. One eligible definition of market impact Δp is the movement of the midpoint, following the placement of an order of size ω . If we define the reference point so that $N\left(a,t\right)\equiv 0$, and the market order is a buy, this definition puts $\omega\left(\Delta p,t\right)=N\left(a+2\Delta p,t\right)-N\left(a,t\right)$. In words, the midpoint shift is half the shift in the best offer. An alternative choice would be to let $\omega\left(\Delta p,t\right)=N\left(\Delta p,t\right)$, which would include part of the instantaneous spread in the definition of impact in midpoint-centered coordinates, or none of it in ask-centered coordinates. The issue of how impact is related to $N\left(p,t\right)$ is separate from whether the best ask is set equal to the reference point for prices, and may be chosen differently to answer different questions.

Under any such definition, however, the impact Δp is a monotonic function of ω in every instance, so either may be taken as the independent variable, along with the index t that labels the instance. We wish to account for the differences in instance averages $\langle \rangle$ of ω and Δp , regarded respectively as the dependent variables, in terms of the fluctuations of either.

In spite of the fact that the density $n\left(p,t\right)$ is a highly discontinuous variable in general, monotonicity of the cumulative $N\left(p,t\right)$ enables us to picture a power series expansion for $\omega\left(p,t\right)$ in p, with coefficients that fluctuate in time. The simplest such expansion that captures much of the behavior of the simulated output is

$$\omega(p,t) = a(t) + b(t)p + \frac{c(t)}{2}p^2,$$
 (A1)

if p is regarded as the independent variable, or

$$p\left(\omega,t\right) = \frac{-b\left(t\right) + \sqrt{b^{2}\left(t\right) + 2c\left(t\right)\left(\omega - a\left(t\right)\right)}}{c\left(t\right)}, \quad (A2)$$

if ω is. While the variable a(t) would seem unnecessary since ω is zero at p=0, empirically we find that

simultaneous fits to both ω and ω^2 at low order can be made better by incorporating the additional freedom of fluctuations in a.

We imagine splitting each t-dependent coefficient into its mean, and a zero-mean fluctuation component, as

$$a(t) \equiv \bar{a} + \delta a(t), \tag{A3}$$

$$b(t) \equiv \bar{b} + \delta b(t) \tag{A4}$$

and

$$c(t) \equiv \bar{c} + \delta c(t). \tag{A5}$$

The fluctuation components will in general depend on ϵ . The values of the mean and second moment of the fluctuations can be extracted from the mean distributions $\langle \omega \rangle$ and $\langle \omega^2 \rangle$. The mean values come from the linear expectation:

$$\langle \omega \left(0 \right) \rangle = \bar{a},\tag{A6}$$

$$\left. \frac{\partial \left\langle \omega \left(p \right) \right\rangle}{\partial p} \right|_{p=0} = \bar{b},\tag{A7}$$

and

$$\left. \frac{\partial^2 \left\langle \omega \left(p \right) \right\rangle}{\partial p^2} \right|_{p=0} = \bar{c}. \tag{A8}$$

Given these, the fluctuations then come from the quadratic expectation as

$$\langle \omega^2(0) \rangle = \bar{a}^2 + \langle \delta a^2 \rangle,$$
 (A9)

$$\left. \frac{\partial \left\langle \omega^2 \left(p \right) \right\rangle}{\partial p} \right|_{p=0} = 2\bar{a}\bar{b} + 2 \left\langle \delta a \delta b \right\rangle, \tag{A10}$$

$$\left. \frac{\partial^2 \left\langle \omega^2 \left(p \right) \right\rangle}{\partial p^2} \right|_{p=0} = 2 \left(\bar{b}^2 + \bar{a}\bar{c} \right) + 2 \left\langle \delta b^2 + \delta a \delta c \right\rangle, \quad (A11)$$

$$\left. \frac{\partial^3 \left\langle \omega^2 \left(p \right) \right\rangle}{\partial p^3} \right|_{p=0} = 6 \left(\bar{b}\bar{c} + \left\langle \delta b \delta c \right\rangle \right), \tag{A12}$$

and

$$\left. \frac{\partial^4 \left\langle \omega^2 \left(p \right) \right\rangle}{\partial p^4} \right|_{p=0} = 6 \left(\bar{c}^2 + \left\langle \delta c^2 \right\rangle \right). \tag{A13}$$

When ω is given a specific definition in terms of the cumulative distribution, its averages become averages over the density in the order book.

The values of the moments as obtained above may then be used in a derivative expansion of the inverse function (A2), making the prediction for the averaged impact

$$\langle p(\omega) \rangle = \bar{p} + \frac{1}{2} \frac{\overline{\partial^2 p}}{\partial a^2} \left\langle \delta a^2 \right\rangle + \frac{1}{2} \frac{\overline{\partial^2 p}}{\partial b^2} \left\langle \delta b^2 \right\rangle + \frac{1}{2} \frac{\overline{\partial^2 p}}{\partial c^2} \left\langle \delta c^2 \right\rangle + \frac{1}{2} \frac{\overline{\partial^2 p}}{\partial c^2} \left\langle \delta a \delta b \right\rangle + \frac{\overline{\partial^2 p}}{\partial a \partial b} \left\langle \delta a \delta b \right\rangle + \frac{\overline{\partial^2 p}}{\partial a \partial c} \left\langle \delta a \delta c \right\rangle + \frac{\overline{\partial^2 p}}{\partial b \partial c} \left\langle \delta b \delta c \right\rangle,$$
(A14)

where overbar denotes the evaluation of the function (A2) or its indicated derivative at $b(t) = \bar{b}$, $c(t) = \bar{c}$, and ω . The fluctuations $\langle \delta b^2 \rangle$ and $\langle \delta a \delta c \rangle$ cannot be determined independently from Eq. (A11). However, in keeping with this fact, their coefficient functions in Eq. (A14) are identical, so the inversion remains fully specified.

If we denote by \bar{Z} the radical

$$\bar{Z} \equiv \sqrt{\bar{b}^2 + 2\bar{c}(\omega - \bar{a})},\tag{A15}$$

the various partial derivative functions in Eq. (A14) evaluate to

$$\frac{1}{2}\frac{\overline{\partial^2 p}}{\partial a^2} = \frac{-\bar{c}}{2\bar{Z}^3},\tag{A16}$$

$$\frac{\overline{\partial^2 p}}{\partial a \partial b} = \frac{\bar{b}}{\bar{Z}^3},\tag{A17}$$

$$\frac{1}{2}\frac{\overline{\partial^2 p}}{\partial h^2} = \overline{\frac{\partial^2 p}{\partial a \partial c}} = \frac{\omega - \bar{a}}{\bar{Z}^3},\tag{A18}$$

$$\frac{\overline{\partial^2 p}}{\partial b \partial c} = \frac{1}{\bar{c}^2} - \frac{\bar{b}}{\bar{Z}\bar{c}^2} - \frac{(\omega - \bar{a})\bar{b}}{\bar{c}\bar{Z}^3}, \tag{A19}$$

and

$$\frac{1}{2}\frac{\overline{\partial^2 p}}{\partial c^2} = \frac{\bar{Z} - \bar{b}}{\bar{c}^3} - \frac{\omega - \bar{a}}{2\bar{c}^2\bar{Z}} - \frac{(\omega - \bar{a})^2}{2\bar{c}\bar{Z}^3}.$$
 (A20)

Plugging these into Eq. (A14) gives the predicted mean price impact, compared to actual mean in Fig. 31. Here the measure used for price impact is the movement of the ask from buy market orders. The cumulative order distribution is computed in ask-centered coordinates, eliminating the contribution from the half-spread in the p coordinate. The inverse of the mean cumulative distribution (dotted), which corresponds to \bar{p} in Eq. (A14), clearly underestimates the actual mean impact (solid). However, the corrections from only second-order fluctuations in a, b, and c account for much of the difference at all values of ϵ .

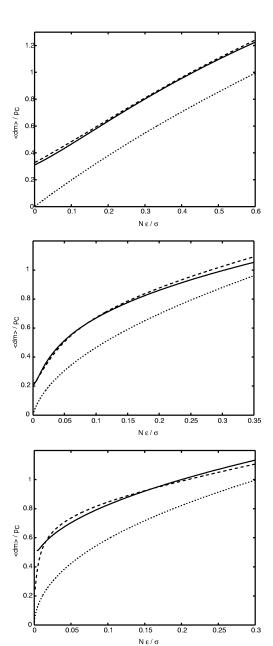


FIG. 31: Comparison of the inverse mean cumulative order distribution \bar{p} (dot), to the actual mean impact (solid), and the second-order fluctuation expansion (A14, dash). (a) : $\epsilon = 0.2$. (b): $\epsilon = 0.02$. (c): $\epsilon = 0.002$.

2. Quantiles

Another way to characterize the relationship between depth profile and market impact is in terms of their quantiles (the fraction greater than a given value, for example the median is the 0.5 quantile). Interestingly, the relationship between quantiles is trivial. Letting $Q_r(x)$ be the r^{th} quantile of x, because the cumulative depth N(p) is a non-decreasing function with inverse $p = \phi(N)$,

we have the relation

$$Q_r(\phi) = (Q_{1-r}(N))^{-1} \tag{A21}$$

This provides an easy and accurate way to compare depth and price impact when the tick size is sufficiently small. However, when the tick size is very coarse, the quantiles are in general not very useful, because unlike the mean, the quantiles do not vary continuously, and only take on a few discrete values.

As we have argued in the previous section, in nondimensional coordinates all of the properties of the limit order book are described by the two dimensionless scale factors ϵ and dp/p_c (see table (II). When expressed in dimensionless coordinates, any property, such as depth, spread, or price impact, can only depend on these two parameters. This reduces the search space from five dimensions to two, which greatly simplifies the analysis. Any results can easily be re-expressed in dimensional coordinates from the definitions of the dimensionless parameters.

APPENDIX B: SUPPORTING CALCULATIONS IN DENSITY COORDINATES

The following two subsections provide details for the master equation solution in density coordinates. The first provides the generating functional solution for the density at general dp, and the second the approximate source term for correlated fluctuations.

1. Generating functional at general bin width

As in the main text, α and μ represent the functions of p everywhere in this subsection, because the boundary values do not propagate globally. Eq. (24) can be solved by assuming there is a convergent expansion in (formal) small D/δ ,

$$\Pi \equiv \sum_{j} \left(\frac{D}{\delta}\right)^{j} \Pi_{j}, \tag{B1}$$

and it is convenient to embellish the shorthand notation as well, with

$$\Pi_i(0, p) \equiv \pi_{0i}(p). \tag{B2}$$

It follows that expected number also expands as

$$\langle n \rangle \equiv \sum_{j} \left(\frac{D}{\delta} \right)^{j} \langle n \rangle_{j}.$$
 (B3)

Order by order in D/δ , Eq. (24) requires

$$\left(\frac{\partial}{\partial \lambda} - \frac{\alpha dp - \mu/2\lambda}{\delta \sigma}\right) \Pi_j = \frac{\mu}{2\delta \sigma \lambda} \pi_{0j} + \frac{\partial}{\partial p^2} \frac{\Pi_{j-1}}{(\lambda - 1)}.$$
(B4)

Because Π_j have been introduced in order to be chosen homogeneous of degree zero in D/δ , the normalization condition requires that

$$\Pi_0(1,p) = 1 , \ \Pi_{i \neq 0}(1,p) = 0 , \ \forall p$$
 (B5)

The implied recursion relations for expected occupation numbers are

$$\langle n \rangle_0 = \frac{\alpha dp}{\delta} - \frac{\mu}{2\delta} \left(1 - \pi_{00} \right), \tag{B6}$$

at j = 0, and

$$\langle n \rangle_j = \frac{\mu}{2\delta} \pi_{0j} + \frac{\partial}{\partial p^2} \langle n \rangle_{j-1}$$
 (B7)

otherwise.

Eq. (B4) is solved immediately by use of an integrating factor, to give the recursive integral relation

$$\Pi_{j}(\lambda) = \pi_{0j} \left[1 + \frac{\alpha dp}{\delta \sigma} \mathcal{I}(\lambda) \right]
+ \mathcal{I}(\lambda) \left\langle \left\langle \frac{\partial}{\partial p^{2}} \frac{\Pi_{j-1}}{\lambda - 1} \right\rangle \right\rangle,$$
(B8)

where

$$\mathcal{I}(\lambda) \equiv \lambda \int_0^1 dz e^{(\alpha dp/\delta \, \sigma)(1-z)} z^{(\mu/2\delta \, \sigma)}, \qquad (B9)$$

and

$$\left\langle \left\langle \frac{\partial}{\partial p^2} \frac{\Pi_{j-1}}{\lambda - 1} \right\rangle \right\rangle_{\lambda} \equiv \frac{\lambda}{\mathcal{I}(\lambda)} \int_{0}^{1} dz \quad e^{(\alpha dp/\delta \, \sigma)(1-z)} z^{(\mu/2\delta \, \sigma)} \frac{\partial^{2}}{\partial p^{2}} \frac{\Pi_{j-1}(\lambda z)}{\lambda z - 1}.$$
(B10)

The surface condition (B5) provides the starting point for this recursion, by giving at j = 0

$$\pi_{00} = \frac{1}{1 + (\alpha dp/\delta \sigma) \mathcal{I}(1)}.$$
 (B11)

Given forms for α and μ , Eq. (B6) may be solved directly from Eq. (B11), and extended by Eq. (B7) to solve for $\langle n(p) \rangle$. More generally, equations (B9), (B10), and (B11) may be solved to any desired order numerically, to obtain the fluctuation characteristics of n(p). Finding the solution becomes difficult, however, when α and μ must be related self-consistently to the solutions for Π . The special case $dp \rightarrow 0$ admits a drastic simplification, in which the whole expansion for $\langle n(p) \rangle$ may be directly summed, to recover the result in the main text. In this limit, one gets a single differential equation in p which is solvable by numerical integration. The existence and regularity of this solution demonstrates the existence of a continuum limit on the price space, and can be simulated directly by allowing orders to be placed at arbitrary real-valued prices.

a. Recovering the continuum limit for prices

In the limit that the dimensionless quantity $\alpha dp/\delta \sigma \rightarrow 0$, Eq. (B9) simplifies to

$$\mathcal{I}\left(\lambda\right)
ightarrow rac{\lambda}{1 + \mu/2\delta\,\sigma} + \mathcal{O}\left(dp
ight),$$
 (B12)

from which it follows that

$$\langle n \rangle_0 \to \frac{\alpha dp/\delta}{1 + \mu/2\delta \sigma} + \mathcal{O}\left(dp^2\right).$$
 (B13)

The important simplification given by vanishing dp, as will be seen below, is that the expansion (19) collapses, at leading order in dp, to

$$\Pi_0(\lambda) \to 1 + (\lambda - 1) \frac{\langle n \rangle_0}{\sigma} + \mathcal{O}(dp^2).$$
 (B14)

Eq. (B14) is used as the input to an inductive hypothesis

$$\Pi_{j-1}(\lambda) \to \Pi_{j-1}(1) + (\lambda - 1) \frac{\langle n \rangle_{j-1}}{\sigma} + \mathcal{O}\left(dp^2\right), \text{ (B15)}$$

(n. b. $\langle n \rangle_{j-1} \sim \mathcal{O}(dp)$, $\Pi_{j-1}(1) = \text{ either 1 or 0}$), which with Eq. (B10), then recovers the condition at j:

$$\Pi_{j}(\lambda) \to (\lambda - 1) \mathcal{I}(1) \frac{d^{2}}{dp^{2}} \frac{\langle n \rangle_{j-1}}{\sigma} + \mathcal{O}(dp^{2}).$$
 (B16)

Using Eq. (19) at $\lambda \to 1$, and Eq. (B12) for \mathcal{I} , gives the recursion for the number density

$$\langle n \rangle_{j \neq 0} \rightarrow \frac{1}{1 + \mu/2\delta \,\sigma} \, \frac{d^2}{dp^2} \langle n \rangle_{j-1} + \mathcal{O}\left(dp^2\right).$$
 (B17)

The sum (B3) for $\langle n \rangle$ is then

$$\langle n \rangle = \sum_{j} \left(\frac{D}{\delta} \frac{1}{1 + \mu/2\delta \sigma} \frac{d^2}{dp^2} \right)^{j} \langle n \rangle_0.$$
 (B18)

Using Eq. (B13) for $\langle n \rangle_0$ and re-arranging terms, Eq. (B18) is equivalent to

$$\langle n \rangle = \frac{1}{1 + \mu/2\delta \,\sigma} \sum_{j} \left(\frac{D}{\delta} \, \frac{d^2}{dp^2} \, \frac{1}{1 + \mu/2\delta \,\sigma} \right)^{j} \frac{\alpha dp}{\delta}. \tag{B19}$$

The series expansion in the price Laplacian is formally the geometric sum

$$(1 + \mu/2\delta\sigma)\langle n \rangle = \left[1 - \frac{D}{\delta} \frac{d^2}{dp^2} \frac{1}{1 + \mu/2\delta\sigma}\right]^{-1} \frac{\alpha dp}{\delta},$$
(B20)

which can be inverted to give Eq. (27), a relation that is local in derivatives.

I	case	source	prob
	$\Delta \hat{p} \leq \hat{a}' < \hat{p}$	$\psi\left(\hat{p} - \Delta\hat{p}\right) - \psi\left(\hat{p}\right)$	$\varphi\left(\Delta\hat{p}\right) - \varphi\left(\hat{p}\right)$
	$\Delta \hat{p} \le \hat{p} < \hat{a}' \le \hat{p} + \Delta \hat{p}$	$0 - \psi\left(\hat{p}\right)$	$\varphi\left(\hat{p}\right) - \varphi\left(\hat{p} + \Delta\hat{p}\right)$
	$\hat{p} < \Delta \hat{p} < \hat{a}' \le \hat{p} + \Delta \hat{p}$	$0 - \psi\left(\hat{p}\right)$	$\varphi\left(\Delta\hat{p}\right) - \varphi\left(\hat{p} + \Delta\hat{p}\right)$

TABLE VII: Contributions to "effective P_- " from removal of a buy limit order, conditioned on the position of the ask relative to p.

case	source	prob
$\Delta \hat{p} \leq \hat{a}' < \hat{p}$	$\psi\left(\hat{p} + \Delta\hat{p}\right) - \psi\left(\hat{p}\right)$	$\varphi\left(\Delta\hat{p}\right) - \varphi\left(\hat{p}\right)$
$\Delta \hat{p} \le \hat{p} < \hat{a}' \le \hat{p} + \Delta \hat{p}$	$\psi\left(\hat{p} + \Delta\hat{p}\right) - 0$	$\varphi\left(\hat{p}\right) - \varphi\left(\hat{p} + \Delta\hat{p}\right)$
$\hat{p} < \Delta \hat{p} < \hat{a}' \le \hat{p} + \Delta \hat{p}$	$\psi\left(\hat{p} + \Delta\hat{p}\right) - 0$	$\varphi\left(\Delta\hat{p}\right) - \varphi\left(\hat{p} + \Delta\hat{p}\right)$

TABLE VIII: Contributions to "effective P_+ " from removal of a sell limit order, conditioned on the position of the ask relative to p.

2. Cataloging correlations

A correct source term \mathcal{S} must correlate the incidences of zero occupation with the events producing shifts. It is convenient to separate these into the four independent types of deposition and removal.

First we consider removal of buy limit orders, which generates a negative shift of the midpoint. Let \hat{a}' denote the position of the ask after the shift. Then all possible shifts $\Delta \hat{p}$ are related to a given price bin \hat{p} and \hat{a}' in one of three ordering cases, shown in Table VII. For each case, the source term corresponding to $[\psi (\hat{p} - \Delta \hat{p}) - \psi (\hat{p})]$ in Eq. (43) is given, together with the measure of order-book configurations for which that case occurs. The mean-field assumption (36) is used to estimate these measures.

As argued when defining β in the simpler diffusion approximation for the source terms, the measure of shifts from removal of either buy or sell limit orders should be symmetric with that of their addition within the spread, which is is $2d\Delta\hat{p}$ for either type, in cases when the shift $\pm\Delta\hat{p}$ is consistent with the value of the spread. The only change in these more detailed source terms is replacement of the simple $\Pr\left(a\geq\Delta p\right)$ with the entries in the third column of Table VII. When the $\Delta\hat{p}$ cases are integrated over their range as specified in the first column and summed, the result is a contribution to \mathcal{S} of

$$\int_{0}^{\hat{p}} 2d\Delta \hat{p} \,\psi\left(\hat{p} - \Delta \hat{p}\right) \left[\varphi\left(\Delta \hat{p}\right) - \varphi\left(\hat{p}\right)\right] \\ - \int_{0}^{\infty} 2d\Delta \hat{p} \,\psi\left(\hat{p}\right) \left[\varphi\left(\Delta \hat{p}\right) - \varphi\left(\hat{p} + \Delta \hat{p}\right)\right] \text{ (B21)}$$

Sell limit-order removals generate another sequence of cases, symmetric with the buys, but inducing positive shifts. The cases, source terms, and frequencies are given in Table VIII. Their contribution to \mathcal{S} , after integration

case	source	prob
$\Delta \hat{p} \le \hat{a} < \hat{p} - \Delta \hat{p}$	$\psi\left(\hat{p} - \Delta\hat{p}\right) - \psi\left(\hat{p}\right)$	$\varphi\left(\Delta\hat{p}\right) - \varphi\left(\hat{p} - \Delta\hat{p}\right)$
$\Delta \hat{p} \le \hat{p} - \Delta \hat{p} < \hat{a} \le \hat{p}$	$0 - \psi\left(\hat{p}\right)$	$\varphi\left(\hat{p} - \Delta\hat{p}\right) - \varphi\left(\hat{p}\right)$
$\hat{p} - \Delta \hat{p} < \Delta \hat{p} < \hat{a} \le \hat{p}$	$0 - \psi\left(\hat{p}\right)$	$\varphi\left(\Delta\hat{p}\right) - \varphi\left(\hat{p}\right)$

TABLE IX: Contributions to "effective P_{-} " from addition of a sell limit order, conditioned on the position of the ask relative to p.

case	source	prob
$\Delta \hat{p} \le \hat{a}' < \hat{p}$	$\psi\left(\hat{p} + \Delta\hat{p}\right) - \psi\left(\hat{p}\right)$	$\varphi\left(\Delta\hat{p}\right) - \varphi\left(\hat{p}\right)$
$\Delta \hat{p} \le \hat{p} < \hat{a}' \le \hat{p} + \Delta \hat{p}$	$\psi\left(\hat{p} + \Delta\hat{p}\right) - 0$	$\varphi\left(\hat{p}\right) - \varphi\left(\hat{p} + \Delta\hat{p}\right)$
$\hat{p} < \Delta \hat{p} < \hat{a}' \le \hat{p} + \Delta \hat{p}$	$\psi\left(\hat{p} + \Delta\hat{p}\right) - 0$	$\varphi\left(\Delta\hat{p}\right) - \varphi\left(\hat{p} + \Delta\hat{p}\right)$

TABLE X: Contributions to "effective P_+ " from addition of a buy limit order, conditioned on the position of the ask relative to p.

over $\Delta \hat{p}$, is then

$$\int_{0}^{\infty} 2d\Delta \hat{p} \,\psi\left(\hat{p} + \Delta \hat{p}\right) \left[\varphi\left(\Delta \hat{p}\right) - \varphi\left(\hat{p} + \Delta \hat{p}\right)\right] - \int_{0}^{\hat{p}} 2d\Delta \hat{p} \,\psi\left(\hat{p}\right) \left[\varphi\left(\Delta \hat{p}\right) - \varphi\left(\hat{p}\right)\right]$$
(B22)

Order addition is treated similarly, except that \hat{a} denotes the position of the ask before the event. Sell limit-order additions generate negative shifts, with the cases shown in Table IX. Integration over $\Delta \hat{p}$ consistent with these cases gives the negative-shift contribution to \mathcal{S}

$$\int_{0}^{\hat{p}/2} 2d\Delta \hat{p} \,\psi\left(\hat{p} + \Delta \hat{p}\right) \left[\varphi\left(\Delta \hat{p}\right) - \varphi\left(\hat{p} - \Delta \hat{p}\right)\right] - \int_{0}^{\hat{p}} 2d\Delta \hat{p} \,\psi\left(\hat{p}\right) \left[\varphi\left(\Delta \hat{p}\right) - \varphi\left(\hat{p} + \Delta \hat{p}\right)\right]$$
(B23)

The corresponding buy limit-order addition cases are given in Table X, and their positive-shift contribution to S turns out to be the same as that from removal of sell limit orders (B22).

Writing the source as a sum of two terms $S \equiv S_{\rm buy} + S_{\rm sell}$, the combined contribution from buy limit-order additions and removals is

$$S_{\text{buy}}(\hat{p}) = \int_{0}^{\hat{p}} 2d\Delta \hat{p} \left[\psi \left(\hat{p} - \Delta \hat{p} \right) - \psi \left(\hat{p} \right) \right] \left[\varphi \left(\Delta \hat{p} \right) - \varphi \left(\hat{p} \right) \right]$$
$$- \int_{0}^{\infty} 2d\Delta \hat{p} \left[\psi \left(\hat{p} + \Delta \hat{p} \right) - \psi \left(\hat{p} \right) \right] \left[\varphi \left(\Delta \hat{p} \right) - \varphi \left(\hat{p} + \Delta \hat{p} \right) \right]$$
(B24)

The corresponding source term from sell order addition

and removal is

$$S_{\text{sell}}(\hat{p}) = \int_{0}^{\hat{p}/2} 2d\Delta \hat{p} \, \psi \, (\hat{p} - \Delta \hat{p}) \left[\varphi \, (\Delta \hat{p}) - \varphi \, (\hat{p} - \Delta \hat{p}) \right]$$
$$- 2 \int_{0}^{\hat{p}} 2d\Delta \hat{p} \, \psi \, (\hat{p}) \left[\varphi \, (\Delta \hat{p}) - \varphi \, (\hat{p}) \right]$$
$$+ \int_{0}^{\infty} 2d\Delta \hat{p} \, \psi \, (\hat{p} + \Delta \hat{p}) \left[\varphi \, (\Delta \hat{p}) - \varphi \, (\hat{p} + \Delta \hat{p}) \right]$$
(B25)

The forms (B24) and (B25) do not lead to $\int d\hat{p}S = 0$, and correcting this presumably requires distributing the orders erroneously transported through the midpoint by the diffusion term, to interior locations where they then influence long-time diffusion autocorrelation. These source terms manifestly satisfy S(0) = 0, though, and that determines the intercept of the average order depth.

a. Getting the intercept right

Evaluating Eq. (44) with $\alpha(\hat{p})/\alpha(\hat{\infty}) = 1 + \Pr(\hat{s}/2 \geq \hat{p})$, at $\hat{p} = 0$ gives the boundary value of the nondimensionalized, midpoint-centered, mean order density

$$\psi\left(0\right) = \frac{2}{1+\epsilon},\tag{B26}$$

which dimensionalizes to

$$\frac{\langle n(0) \rangle}{\sigma dp} = \frac{2\alpha(\infty)/\sigma}{\mu(0)/2\sigma + \delta}.$$
 (B27)

Eq. (B27), for the *total* density, is the same as the form (B6) produced by the diffusion solution for the *zeroth order* density, as should be the case if diffusion no longer transports orders through the midpoint. This form is verified in simulations, with midpoint-centered averaging.

Interestingly, the same argument for the bid-centered frame would simply omit the φ from $\alpha(0)/\alpha(\infty)$, predicting that

$$\psi\left(0\right) = \frac{1}{1+\epsilon},\tag{B28}$$

a result which is not confirmed in simulations. Thus, in addition to not satisfying the mean-field approximation, the bid-centered density average appears to receive some diffusive transport of orders all the way down to the bid.

b. Fokker-Planck expanding correlations

Equations (B24) and (B25) are not directly easy to use in a numerical integral. However, they can be Fokker-Planck expanded to terms with behavior comparable to the diffusion equation, and the correct behavior near the midpoint. Doing so gives the nondimensional expansion of the source term S corresponding to the diffusion contribution in Eq. (32):

$$S = \mathcal{R}(\hat{p}) \psi(\hat{p}) + \mathcal{P}(\hat{p}) \frac{d\psi(\hat{p})}{d\hat{p}} + \epsilon \beta(\hat{p}) \frac{d^2 \psi(\hat{p})}{d\hat{p}^2}. \quad (B29)$$

The rate terms in Eq. (B29) are integrals defined as

$$\mathcal{R}(\hat{p}) = \int_{0}^{\infty} 2d\Delta \hat{p} \left[\varphi \left(\Delta \hat{p} \right) - \varphi \left(\hat{p} + \Delta \hat{p} \right) \right]$$
$$-2 \int_{0}^{\hat{p}} 2d\Delta \hat{p} \left[\varphi \left(\Delta \hat{p} \right) - \varphi \left(\hat{p} \right) \right]$$
$$+ \int_{0}^{\hat{p}/2} 2d\Delta \hat{p} \left[\varphi \left(\Delta \hat{p} \right) - \varphi \left(\hat{p} - \Delta \hat{p} \right) \right] (B30)$$

and

$$\epsilon\beta\left(\hat{p}\right) = \int_{0}^{\infty} 2d\Delta\hat{p} \left(\Delta\hat{p}\right)^{2} \left[\varphi\left(\Delta\hat{p}\right) - \varphi\left(\hat{p} + \Delta\hat{p}\right)\right]$$

$$+ \frac{1}{2} \int_{0}^{\hat{p}} \left(\Delta\hat{p}\right)^{2} 2d\Delta\hat{p} \left[\varphi\left(\Delta\hat{p}\right) - \varphi\left(\hat{p}\right)\right]$$

$$+ \frac{1}{2} \int_{0}^{\hat{p}/2} 2d\Delta\hat{p} \left(\Delta\hat{p}\right)^{2} \left[\varphi\left(\Delta\hat{p}\right) - \varphi\left(\hat{p} - \Delta\hat{p}\right)\right].$$
(B32)

All of the coefficients (B30 - B31) vanish manifestly as $\hat{p} \to 0$, and at large \hat{p} , $\mathcal{R}, \mathcal{P} \to 0$, while $\epsilon \beta(\hat{p}) \to 4 \int_0^\infty d\Delta \hat{p} \left(\Delta \hat{p}\right)^2 \varphi(\Delta \hat{p})$, recovering the diffusion constant (41) of the simplified source term. However, they are still not convenient for numerical integration, being nonlocal in φ .

The exponential form (33) is therefore exploited to approximate φ , in the region where its value is largest, with

the expansion

$$\varphi(\hat{p} \pm \Delta \hat{p}) \approx \varphi(\hat{p}) \varphi(\Delta \hat{p}) e^{\pm \hat{p} \Delta \hat{p} \partial \psi/\partial \hat{p}|_{0}}$$
 (B33)

In the range where the mean-field approximation is valid, φ is dominated by the constant term ψ (0), and even the factors $e^{\pm\hat{p}\,\Delta\hat{p}\,\partial\psi/\partial\hat{p}|_0}$ can be approximated as unity. This leaves the much-simplified expansions

$$\mathcal{R}(\hat{p}) = [1 - \varphi(\hat{p})] \mathcal{I}_{0}(\infty) - 2 [\mathcal{I}_{0}(\hat{p}) - 2\hat{p}\varphi(\hat{p})] + [1 - \varphi(\hat{p})] \mathcal{I}_{0}(\hat{p}/2),$$
(B34)

$$\mathcal{P}(\hat{p}) = 2 \left[1 - \varphi(\hat{p}) \right] \mathcal{I}_{1}(\infty) - \left[\mathcal{I}_{1}(\hat{p}) - \hat{p}^{2} \varphi(\hat{p}) \right] - \left[1 - \varphi(\hat{p}) \right] \mathcal{I}_{1}(\hat{p}/2),$$
(B35)

and

$$\epsilon \beta \left(\hat{p} \right) = \left[1 - \varphi \left(\hat{p} \right) \right] \mathcal{I}_{2} \left(\infty \right)$$

$$+ \frac{1}{2} \left[\mathcal{I}_{2} \left(\hat{p} \right) - \frac{2}{3} \hat{p}^{3} \varphi \left(\hat{p} \right) \right]$$

$$+ \left[1 - \varphi \left(\hat{p} \right) \right] \mathcal{I}_{2} \left(\hat{p} / 2 \right).$$
 (B36)

In Equations (B34 - B36),

$$\mathcal{I}_{j}\left(\hat{p}\right) \equiv \int_{0}^{\hat{p}} 2d\Delta \hat{p} \left(\Delta \hat{p}\right)^{j} \varphi\left(\Delta \hat{p}\right), \tag{B37}$$

for j=0,1,2. These forms (B34 - B36) are inserted in Eq. (B29) for $\mathcal S$ to produce the mean-field results compared to simulations in Fig. 20 - Fig. 22.

ACKNOWLEDGMENTS

We would like to thank the McKinsey Corporation, Credit Suisse First Boston, the McDonnel Foundation, Bob Maxfield, and Bill Miller for supporting this research. We would also like to thank Paolo Patelli, R. Rajesh, Spyros Skouras, and Ilija Zovko for helpful discussions, and Marcus Daniels for valuable technical support

^[1] M. Daniels, J.D. Farmer, G. Iori, and D. E. Smith, "How storing supply and demand affects price diffusion", http://xxx.lanl.gov/cond-mat/0112422 (2001).

^[2] L. Bachelier, "Théorie de la spéculation", 1900. Reprinted in P.H. Cootner, The Random Character of Stock Prices, 1964, MIT Press Cambridge.

^[3] D. Gode and S. Sunder, "Allocative efficiency of markets with zero intelligence traders: Markets as a partial substitute for individual rationality", J. of Political Economy, 101:119 (1993).

^[4] J.D. Farmer, "Market force, ecology, and evolution", SFI

working paper 98-12-117 or $\label{eq:http://xxx.lanl.gov/adaptorg/9812005} \text{ and } \text{paper 98-12-117 or } \text{http://xxx.lanl.gov/adaptorg/9812005}.$

^[5] J.-P. Bouchaud and R. Cont, "A Langevin approach to stock market fluctuations and crashes", European Physics Journal B 6, 543 (1998)

^[6] J.A. Hausman and A.W. Lo, "An ordered probit analysis of transaction stock prices", *Journal of Financial Economics* 31 319-379 (1992)

^[7] J.D. Farmer, "Slippage 1996", Prediction Company internal technical report (1996), http://www.predict.com

^[8] N. Torre, BARRA Market Impact Model Handbook,

- BARRA Inc, Berkeley CA, www.barra.com (1997).
- [9] A. Kempf and O. Korn, "Market depth and order size", University of Mannheim technical report (1998).
- [10] V. Plerou, P. Gopikrishnan, X. Gabaix, and H.E. Stanley, "Quantifying stock price response to demand fluctuations", http://xxx.lanl.gov/cond-mat/0106657.
- [11] F. Lillo, J.D. Farmer, and R. Mantegna, "Single curve collapse of the price impact function", to appear in Nature. A more detailed version can be found at http://xxx.lanl.gov/cond-mat/0207438.
- [12] H. Mendelson, "Market behavior in a clearing house", Econometrica 50, 1505-1524 (1982).
- [13] K.J. Cohen, R.M. Conroy and S.F. Maier, "Order flow and the quality of the market", in: Y. Amihud, T. Ho and R. Schwartz, eds., Market Making and the Changing Structure of the Securities Industry, 1985, Lexington Books, Lexington MA.
- [14] I. Domowitz and Jianxin Wang, "Auctions as algorithms", J. of Econ. Dynamics and Control 18, 29 (1994).
- [15] T. Bollerslev, I. Domowitz, and J. Wang, "Order flow and the bid-ask spread: An empirical probability model of screen-based trading", J. of Econ. Dynamics and Control 21, 1471 (1997).
- [16] P. Bak, M. Paczuski, and M. Shubik, "Price variations in a stock market with many agents", *Physica A* 246, 430 (1997).
- [17] D. Eliezer and I.I. Kogan, "Scaling laws for the market microstructure of the interdealer broker markets", http://xxx.lanl.gov/cond-mat/9808240.
- [18] L.-H. Tang and G.-S. Tian, "Reaction-diffusion-

- branching models of stock price fluctuations", *Physica* A, 264, 543 (1999).
- [19] S. Maslov, "Simple model of a limit order-driven market", Physica A 278, 571(2000).
- [20] F. Slanina, "Mean-field approximation for a limit order driven market model", http://xxx.lanl.gov/condmat/0104547
- [21] D. Challet and R. Stinchcombe, "Analyzing and Modeling 1+1d markets", Physica A 300, 285, (2001), preprint cond-mat/0106114
- [22] C. Chiarella and G. Iori, "A simulation analysis of the microstructure of double auction markets", to appear in *Quantitative Finance*, October (2002).
- [23] J. P. Bouchaud, M. Mézard, and M. Potters, "Statistical properties of stock order books: empirical results and models", http://www.arxiv.org/cond-mat/0203511
- [24] P. W. Bridgman, Dimensional Analysis, 1922, Yale University Press, New Haven.
- [25] D. Challet and R. Stinchcombe, "Exclusion particle models of limit order financial markets", http://xxx.lanl.gov/cond-mat/0106114
- [26] I. I. Zovko and J. D. Farmer, "The power of patience: A behavioral regularity in limit order placement" condmat 0206280, to appear in Quantitative Finance, October (2002)
- [27] S. Maslov and M. Mills, "Price fluctuations from the order book perspective — empirical facts and a simple model", Physica A 299, 234 (2001).